

An Empirical Analysis of the Internet Engineering Task Force with Computational Methods

Mladen Karan¹, Prashant Khare¹, Ravi Shekhar², Matthew Russell Barnes¹, Waleed Iqbal¹, Junaid Qadir³, Stephen McQuistin⁴, Richard G. Clegg¹, Patrick Healey¹, Matthew Purver⁵, Ignacio Castro¹, Gareth Tyson⁶, and Colin Perkins⁷

¹Queen Mary University of London

²University of Essex

³Qatar University

⁴University of St Andrews

⁵Jožef Stefan Institute

⁶Hong Kong University of Science & Technology

⁷University of Glasgow

Working Draft – 7 October 2024

Abstract

Technical operation of the Internet requires interoperability between networks, systems, and applications from different vendors as well as cooperation among a growing number of stakeholders. The Internet Engineering Task Force (IETF) is an open standards developing organisation that plays a critical role in enabling this cooperation and interoperability by bringing together interested parties to develop the protocols that power the Internet, such as IP or HTTP. In this paper, we review more 20 years of data about the IETF, studying who participates in the IETF, the emails they exchange, and the documents they produce. We show how these data can be used to understand who develops technical standards for the Internet, and their interests, and how Internet's growth and maturity has given rise to a longer and complex standards development process. We explore where participants gain and exercise influence, and how this is reflected in the language they use and in their interactions.

This paper is a summary and collation of our prior published work, that can be found in full in: [66, 52, 43, 65, 53, 50, 44, 6].

1 Introduction

Network protocol standards are crucial for successful operation of the Internet. A successful standard provides a basis for interoperability between systems developed by competing vendors, and supports the growth of an open ecosystem of products and services. Further, the process by which network protocol standards are developed, comprising multiple rounds of open feedback and review, has proven remarkably effective in designing high-quality and robust protocols, many of which see widespread deployment

and use. Understanding the Internet standards development process, and how it produces successful protocols is, therefore, important if we are to understand the Internet and how it has evolved.

One of the main organisations that develops protocol standards is the Internet Engineering Task Force (IETF). The IETF was founded in 1986, following on from the US Government-funded effort that developed the early Internet. It has since grown to become a global community of network protocol designers, vendors, network operators, and researchers that develop and publish open network protocol standards and operational guidelines. The IETF publishes its standards, and other documents, in the RFC series.¹ This series comprises around 9,000 documents, authored over more than 50 years, and provides a rich history of the development of the Internet and its protocols [35]. This history is further augmented by the rich archive of supporting documents, work-in-progress drafts, email discussion, and other metadata about the standards development process that the IETF makes available.

The availability of this data gives a unique opportunity to explore the development of the Internet. Standards development is an inherently social and political process [17, 81], requiring cooperation and consensus among a growing number of stakeholders. Due to the importance of the Internet, it's critical to gain a coherent understanding of the activities that take place within the IETF, to understand who develops Internet standards and what factors may predict the success of an idea or proposal for a new standard.

In this paper, we explore the operation of the IETF to understand who develops Internet standards, how the participants interact, and what factors may lead to success:

- Section 2 introduces the IETF and the Internet Standards development process, and reviews the available data.

¹<https://www.rfc-editor.org>

- Section 3 considers trends in the production, complexity, and correctness of IETF standards documents over time.
- Section 4 moves on to consider the people who are involved in Internet standards development and how the demographics of the community have shifted over time.
- Section 5 studies the interaction between participants, explores who has influence and impact in the IETF and how that influence is evident in their activities, and the dynamics of communications within the IETF.
- Section 6 considers communications within the IETF from the perspective of language to explore whether there are differences in linguistic patterns and language use in more influential participants, and to consider how the use of language reflects the consensus-driven nature of the standards development process.
- Section 7 discusses what factors lead to success, exploring what makes a document likely to be published as an RFC and what makes a successful author.

Finally, we review related work in Section 8 and draw some conclusions in Section 9. This paper is a summary and collation of our prior published work [66, 52, 43, 65, 53, 50, 44, 6] that is submitted to stimulate discussion.

2 Background and Datasets

We start by presenting an overview of the IETF, and the publication process for RFCs, before outlining the data sources we use within this paper. We also highlight the ethical considerations of accessing and processing this data.

2.1 An IETF Primer

The IETF is an open standards organisation, which develops Internet standards via contributions and collaborations across a number of voluntary stakeholders, including academics, consultants, industry representatives, governments, and civil society organisations. Through extensive collaboration across contributors, the IETF, and associated organisations, develops Internet standards and other documents. These are published by the RFC Editor (<https://www.rfc-editor.org>) in four *publication streams*: the IETF stream, the Internet Research Task Force (IRTF) stream, the Internet Architecture Board (IAB) stream, and the Independent Submission stream.

There is also a fifth, legacy stream, comprising RFCs published prior to the adoption of separate publication streams in July 2007 [59]. While the IETF is an open standards forum that develops technical standards and operational guidelines for the Internet, the IRTF is an associated organisation that promotes longer-term research, and the IAB provides long-range technical direction for Internet development. The Independent Submission stream “allows RFC publication for some documents that are outside the official IETF/IAB/IRTF process but are relevant to the Internet community” (<https://www.rfc-editor.org/about/independent>).

The standards development process is an inherently collaborative activity. Most day-to-day work is conducted on public mailing lists, in conjunction with three plenary meetings and numerous interim working group meetings per year. The mailing lists are broadly split into three categories: announcement lists, where replies are not allowed; non-working group lists, for discussing topics that do not relate to the work within an IETF working group or IRTF research group; and working group and area lists, where technical discussions take place.

The process of RFC publication begins with the submission of an *Internet-Draft*. Whereas anyone can post a draft, not all drafts become RFCs. After a draft is first posted, multiple revisions might take place resulting in multiple versions of the draft. Each new draft is announced on one or more mailing lists related to the topic of the draft, soliciting feedback and encouraging discussion. Drafts are initially posted by individuals. For publication under the IETF stream, drafts must then be adopted by a working group, where, via further revision, the draft may ultimately be published as an RFC. The process of managing drafts, and the degree and type of peer review conducted prior to their submission to the RFC Editor for publication, differs between streams.

Once the technical development of the draft is complete its publication is managed by the RFC Editor, who maintains the master archive of the RFC documents, along with an index of metadata pertaining to the RFCs and their authors. Finally, once an RFC has been published, deployment is voluntary, and therefore not all RFCs are widely implemented.

2.2 Data Sources

We rely on a number of data sources to conduct our study.

IETF Datatracker: The Datatracker² is the main administrative database used by the IETF, IRTF, and IAB to coordinate their work. It contains information about contributors, working groups, documents, and meetings. The Datatracker provides comprehensive metadata about authors, and the evolution of drafts as they work their way through the standardisation process. The Datatracker was introduced in the early 2000s and its functionality has gradually been extended over time. It contains limited historical data about RFCs produced before its creation, but the IETF has back-filled some data from other records. We have extracted relevant metadata from the Datatracker, using its programmatic REST API, for all RFCs published since 2001. This gives us data pertaining to 4,512 authors, as well as richer metadata for 5,707 RFCs.

RFC Editor: The RFC Editor maintains an index of all RFC publications, alongside metadata related to each document, including its standardisation status, publication stream, authors, and errata. We gather all entries for RFCs published through the end of 2020, giving a total of 8,711 RFCs.

RFCs are not changed after publication, but the RFC Editor maintains a public database of reported errata. The errata filing are separated into *technical errata*, mistakes in the technical content of the RFC that are likely to result in incorrect, non-conforming, implementations of the standards, and *editorial errata* including

²<https://datatracker.ietf.org/>

spelling and punctuation errors that do not otherwise impact the technical content. Errata filed against RFCs published on the IETF stream are checked for correctness by the RFC authors, working group chairs, and area directors. IRTF errata are checked by the authors and the Internet Research Steering Group. The IAB and Independent Submissions Editor check errata for their streams. We collected³ the 6,759 errata reports submitted from January 2001 to the end of December 2022 for analysis.

Email archives and entity resolution: The IETF maintains reasonably complete email archives relating to working group discussions, meeting, and other activities.⁴ We gather all available messages contained within this archive, covering 2,439,240 messages, sent from 74,646 unique email addresses, across 1,153 lists. This snapshot was taken on April 18th 2021.

A key challenge is attributing each email to an individual contributor. Although each, naturally, contains a *From:* header, we find that some contributors use multiple addresses. Beyond this, it is necessary to map each email to the respective person in the Datatracker and RFC Editor datasets. Thus, we perform entity resolution on email senders, and match email senders with their Datatracker profiles. We assign each sender a unique identifier, which is associated with a set of name and address variations from the Datatracker.

Entity resolution takes place in multiple stages. First, we check if the sender's email address has a Datatracker profile. If so, we associate all of their Datatracker metadata with that person's unique identifier, and label the message as having been sent by that identifier. Next, if an email's sender does not appear in the Datatracker, we check, based on their name, if they have already been assigned a Datatracker record. If so, the message is labelled as having been sent by that identifier, and the set of names and addresses associated with that identifier is updated to include the email's sender name and address.

These two stages – matching with the Datatracker, and merging previously seen names and addresses – accounts for the majority (60%) of messages. If an email's sender is not in the Datatracker, and the name and address has not been previously seen, a new person identifier is generated. This accounts for a small portion (10%) of messages. This is reasonable, since Datatracker profiles are necessary for many of the IETF's day-to-day activities.

As a final processing step, we label each identifier as either representing a *contributor*, a *role-based* address, or an *automated* sender. A contributor refers to any standard user participating in the IETF; role-based identifiers reflect addresses that are used by the holders of a particular organisational role, such as the IETF chair; and automated identifiers are those addresses that are system-specific, such as GitHub or IETF notification and announcement addresses. Role-based and automated identifiers account for the remaining 30% of messages.

Note, within the mail archive, spam-indicating headers are present for most messages since 2009, and we confirm pre-filtering is performed by the IETF mail servers. As an extra validation step,

³RFC errata are available at <https://www.rfc-editor.org/errata.php>. We thank the RFC Editor for making the underlying database available to us in machine-readable form for this analysis.

⁴<https://mailarchive.ietf.org/> with public IMAP access also available.

we ran a spam filter over all the archived messages and discarded those marked as spam. Both sources indicate there is very little spam (less than 1%), so it should not have significant impact on our findings.

Reproducibility and data access: We have developed a library, `ietfdata`, that fetches the RFC Index, communicates with the IETF Datatracker API, and retrieves messages from the IETF IMAP mail archives. This library appropriately regulates access, caches data to minimise the impact on the infrastructure, and performs necessary post-processing. This library is available as open source and ongoing development is coordinated on GitHub.⁵

2.3 Ethical Considerations

The data we analyse is extracted from public IETF archives and APIs. We have taken steps to ensure ethical compliance. To ensure that our access to these services does not cause operational problems for the IETF, we are in regular contact with the IETF Tools Team and Secretariat, as well as the operators of the Datatracker and mailing list archive. We have discussed our work with IETF leadership (IETF and IAB Chairs, the IETF Executive Director) and the RFC Editor to ensure that our access falls within their acceptable use policies.

Participants in the IETF must agree to abide by the policies and procedures described at <https://www.ietf.org/about/note-well>, including the privacy policy at <https://www.ietf.org/privacy-statement>. These make explicit provision that mailing list archives and the metadata contained in the Datatracker system will be made public, and it is this public data that we process to extract the aggregate statistics presented in Sections 3, 4, 5, 6 as well as the features analysed in Section 7. We transfer and store data securely, and retain it only for the time needed to perform the analysis. Since we operate entirely using the public APIs provided by IETF, we have no access to private data about individuals.

In balancing these ethical considerations with the reproducibility of our work, we provide the tools needed to access the datasets from the relevant IETF sources, rather than the data itself.

3 Trends in RFC Publication

To begin, we review trends in IETF standards publication. We consider the rate of publication of documents in the RFC series, factors affecting the publication rate, and the occurrence of errata reports, and answer the following research questions:

1. What factors affect the publication rate of RFCs and is the IETF standards development process getting faster or slower over time? (§3.1)
2. Are the standards produced by the IETF broadly correct? Do the reported errata indicate problems with the standards development process in general, or in particular areas, and is the IETF getting better or worse at producing correct RFCs over time? (§3.2)

⁵<https://github.com/glasgow-ip1/ietfdata>

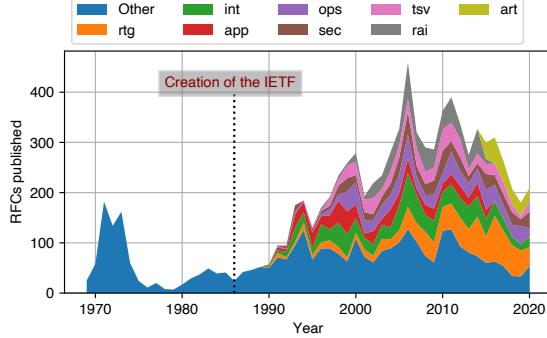


Figure 1: Number of RFCs published over time, subdivided by IETF Area. The areas are *art* – applications and real-time; *app* – applications; *int* – internetworking; *ops* – operations and management; *rai* – real-time applications and infrastructure; *rtg* – routing; *sec* – security; and *tsv* – transport and services. The *other* category includes legacy RFCs, pre-dating the creation of the IETF, and RFCs published in non-IETF publication streams such as the IRTF, IAB, and the Independent stream.

3.1 RFC Production and Complexity

RFC production trends show clear evidence that the IETF is a mature standards development organisation. The peak of activity was in the 1990s and early 2000s, with the initial broad deployment of the Internet, and the organisation has since entered more of a development and maintenance mode of operation. There are clear signs that the complexity of the installed base of protocols is slowing development of new technologies.

RFC Publication Rate: In total, 8,711 RFCs have been published through to the end of 2020. Figure 1 shows how publication trends, in terms of IETF areas and non-IETF streams, have changed over time. We identify three broad publication phases in the RFC series. First, in 1969 through 1974, RFCs are published at a rapid rate during the initial development of the ARPANET. Then, from 1975 through 1985, development slows. This reflects a relatively small community that is gaining real-world experience with the network and developing a small core of applications and protocols. Finally, with the creation of the IETF and the introduction of the National Science Foundation Network (NSFNET) in 1986, both the community, and the number of RFCs published, starts to expand rapidly. This is further driven by the opening of the network to commercial and public use in the mid-1990s. The annual RFC publication rate was highest in 2005, at the peak of the standardisation efforts for SIP and related standards for voice-over-IP and Internet telephony. The rate of publication has slowed in recent years, following the completion of large work programmes relating to HTTP/2 and WebRTC.

Role of Working Groups: From the creation of the IETF in 1986 the growing community, with its interests in an increasing set of applications and protocols, has been split into a number of *working groups*.⁶ These working groups are each chartered with a

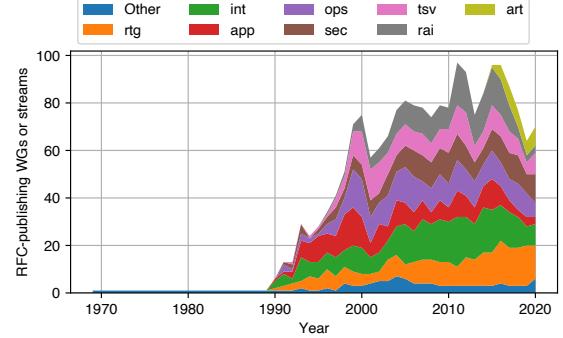


Figure 2: Number of groups publishing RFCs over time, subdivided by IETF Area. The areas are as described in Figure 1.

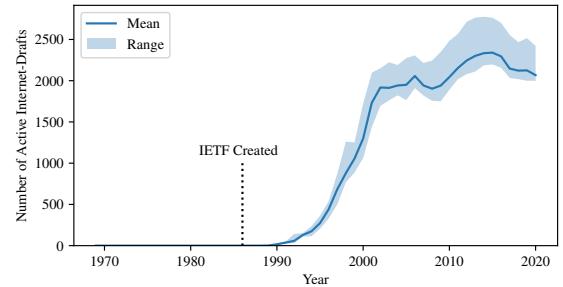


Figure 3: Number of active drafts over time. Internet-Drafts, the working draft series of the IETF, were introduced at the 13th IETF meeting in 1989.

well-defined programme of work and exist within *areas* that have a broader focus.

As shown in Figure 1, the output of different areas has remained relatively stable over time. The most notable trends begin with the creation of the Real-time Applications and Infrastructure (rai) area from within the Transport (tsv) area, and its later merger with the Applications (app) area to become the Applications and Real-Time (art) area around 2014. Additionally, we also observe the significant growth in output of the Routing (rtg) area, owing to the ongoing development of standards for MPLS, service function chaining, and fat tree routing in data centres.

To give a sense for the broader productivity of the IETF, Figure 2 shows the number of working groups that publish RFCs each year. This highlights how the structure of the IETF has grown to accommodate its larger community: in the early 1990s, fewer than 20 working groups were published an RFC each year, while in recent years there has typically been at least 60 working group publishing RFCs, with a peak of 97 working groups publishing RFCs in 2011 (the number of groups publishing RFCs is less than the total number of groups, because it takes time for the work of a new group to reach maturity).

Internet-Drafts: Figure 3 shows the number of working documents, known as Internet-drafts, that were active in the IETF over time. An Internet-draft is a limited lifetime working document that expires after six months, when replaced by a new version, or then published as an RFC.

⁶There are 122 active working groups at the time of writing, organised into 7 areas. The number of groups varies over time, but is typically in the range 100-150.

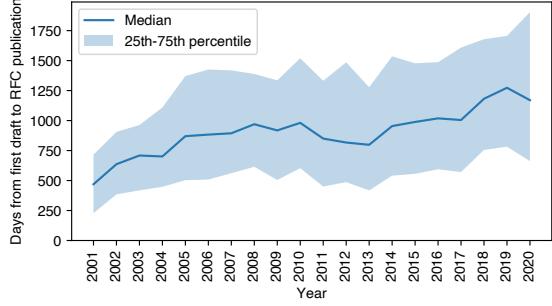


Figure 4: Days from first draft to RFC publication

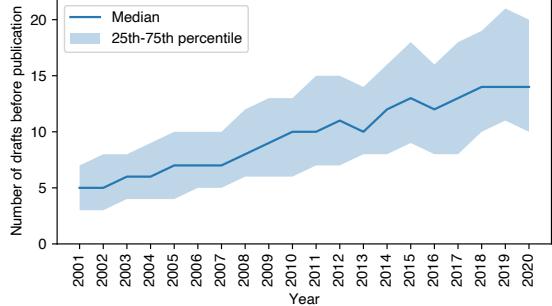


Figure 5: Number of drafts per RFC

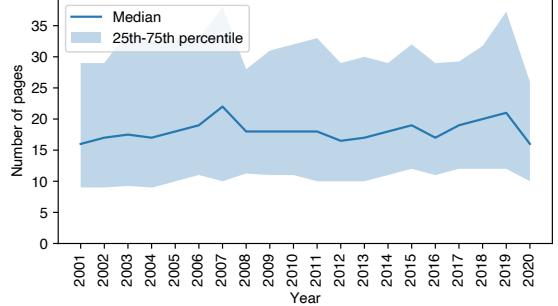


Figure 6: RFC page counts

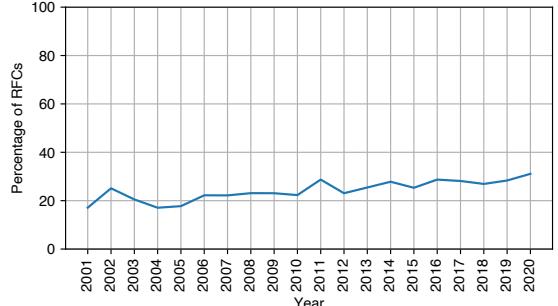


Figure 7: RFCs that update or obsolete previous RFCs

We observe that the number of active drafts correlates more strongly with the number of active working groups (Figure 2) than with RFC publication (Figure 1). This is largely to be expected: an Internet-Draft is a work-in-progress document that is proposed for adoption by, and discussion in, an IETF working group prior to eventual publication as an RFC. Competing proposals are often submitted as drafts, to allow a working group to review them and make a decision how to proceed, and hence not all drafts become RFCs. As we discuss in §7.2, adoption is a significant milestone.

Factors Affecting Publication Rate: The standards development process is taking longer. Figure 4 plots the median number of days from the submission of the first draft of a document through to its publication as an RFC. This shows a clear trend: RFCs are taking longer to make their way through the standardisation and publication process. The median number of days to publication was 469 in 2001, rising to 1,170 in 2020.

Similarly, Figure 5 shows the median number of revisions each document undergoes before being published as an RFC. Days to publication and number of revisions are strongly correlated, suggesting that the time is spent making changes to the document. This may go some way towards explaining the decline in output of the IETF: each RFC is taking longer to produce, with more revisions before publication.

What factors affect the publication rate? One may conjecture that this slowdown is driven by longer RFCs that contain more material. Yet from the median page count of RFCs, shown in Figure 6, no clear trend being visible. The increase in the duration of the standardisation process for RFCs cannot be attributed to RFCs becoming longer: page counts have remained stable.

The publication rate might also be slowing because RFCs are becoming more complex. This could occur, for example, as a result of having to maintain compatibility with older RFCs as the Internet has evolved and matured. Figure 7 shows that this is the case, plotting the proportion of RFCs that are published each year that update (i.e., extend or augment) or obsolete (i.e., replace) one or more previously published RFCs. This percentage has slowly increased as the IETF has matured: in 2020, more than 30% of RFCs updated or made obsolete a previous RFC. Figure 8 expands on this, showing the median number of citations from each RFC to other Internet-Drafts and RFCs. This similarly shows that RFCs are increasingly referring to prior work.

Figure 9 further confirms the growing complexity of RFCs, showing how the use of normative language has evolved. Keywords are used in RFCs to indicate the normative requirements an RFC imposes on implementations [18]. Figure 9 shows the total number of occurrences of each of the keywords (i.e., MUST, MUST NOT, REQUIRED, SHALL, SHALL NOT, SHOULD, SHOULD NOT, RECOMMENDED, MAY, OPTIONAL), divided by the page count of the RFC. As shown, the median number of keywords per page grew from 2001 through to 2010, indicating a growing number of requirements being, before plateauing in recent years.

Summary: RFCs are taking longer to produce, and they go through a greater number of revisions before publication. Further, they increasingly update or reference previously published RFCs, and make greater use of requirements-setting keywords. These are all hallmarks of a maturing standards development organisation that must increasingly consider compatibility with the installed base in the development of its standards.

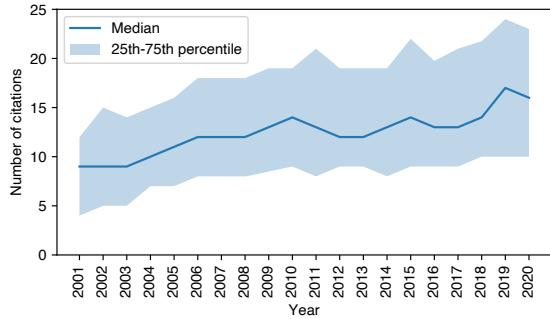


Figure 8: Number of citations to other Internet-Drafts and RFCs per RFC in each year.

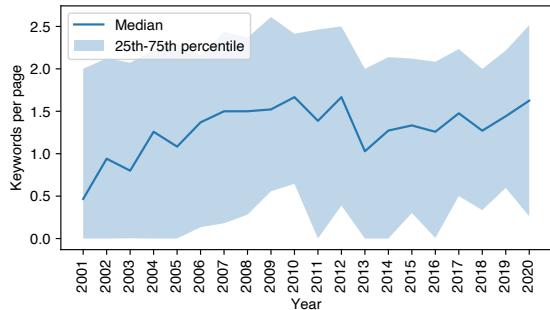


Figure 9: Number of occurrences of normative language keywords per page in RFCs in each year.

3.2 RFC Errata

Mistakes can and do occur in published RFCs. These errors can be reported to the RFC Editor, and the RFC Editor makes the reports publicly available and coordinates with the IETF, and the other publication streams, on their verification. These errata can clarify editorial concerns as well as correct substantive technical errors. Since the presence of significant errors in its specification can undermine the success of a protocol, studying the nature of these errata and their subsequent fixes is important.

In the following, we analyse the 6,759 errata filed with the RFC Editor between 2001–2022, inclusive, documenting 3,288 editorial issues and 3,471 technical issues, and covering 2,240 RFCs, to understand the frequency and causes of problems in RFCs.

Errata over Time: Figure 10 presents the number of errata filed, on average per RFC, since 1969 based on the year of RFC publication. The peak in the number of errata per RFC occurs in 1981. Only 29 RFCs were published that year, but they include major documents such as RFCs 791, 792, and 793 (the original versions of the IP [77], ICMP [76], and TCP [78] standards), with 17, 7, and 47 errata, respectively. These important protocols clearly garnered a great deal of scrutiny and revision. The second highest peak occurs in 2006. In contrast to the previous examples, this has the highest number of RFCs published per year (459), including RFC 4601 [33] that has the most errata (114). Since this second peak, there has been a steady decrease in the number of errata filed. This broadly correlates with the number of RFCs published per year,

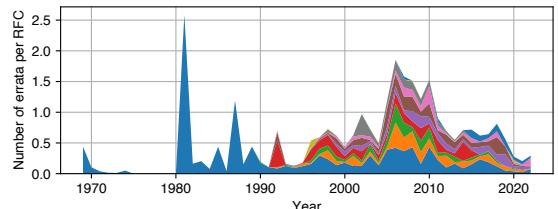


Figure 10: Average number of errata filed per RFC in each year, by RFC publication year, grouped by IETF area. The areas are as described in Figure 1.

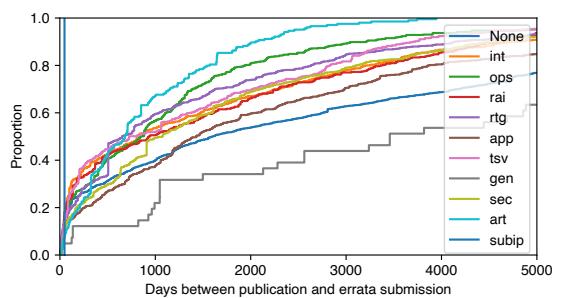


Figure 11: Cumulative distribution of the number of days from RFC publication to errata filing by IETF area. The areas are as described in Figure 1.

with Pearson coefficient 0.59 since 2007. Table 1 lists the top RFCs by errata filing count.

Errata Delay: We next explore how long it takes for errata to be identified and filed. Figure 11 presents a CDF of the number of days between RFC publication and the errata being filed, broken down based on IETF area. We see a wide range of delays. 7.3% of errata are filed within the first 30 days, suggesting that many RFCs are published with issues that could have been identified prior to publication. RFCs from the General (*gen*) area—describing IETF policies and procedures—have the longest delay, with a median of 3,458 days, compared to the Applications and Real-time (*art*) area with a median of 681 days.⁷ Editorial errata are typically filed more quickly, with a median of 987 days, compared to a median of 1,138 days for technical errata.

Errata Status: Figure 12 categorises the errata by verification status and publication year of the RFC to which they relate. The largest share (42.5%) of errata are *verified*: errata that have been confirmed as necessary and accurate. This suggests that many errata identify real problems with RFCs and are useful to the community.

The next largest share (30.3%) are those labelled *hold for document update*. These are errata that are not a necessary update to the RFC, but may be considered on future revisions. For example, *erratum* 6278 describes an oversight in RFC 8610 [12]; the solution to this is non-trivial, and so will be considered in the next version of the specification. Of the 930 RFCs that have *hold for document*

⁷Errata are filed against RFCs within the *subip* area within a median of 48 days, but this is skewed, with only 19 RFCs being published in that area.

| RFC | Title | Year | Area | Filing count |
|------|---|------|------|--------------|
| 4601 | Protocol Independent Multicast - Sparse Mode (PIM-SM): Protocol Specification (Revised) | 2006 | rtg | 114 |
| 4880 | OpenPGP Message Format | 2007 | sec | 52 |
| 793 | Transmission Control Protocol | 1981 | None | 47 |
| 4634 | US Secure Hash Algorithms (SHA and HMAC-SHA) | 2006 | None | 44 |
| 5661 | Network File System (NFS) Version 4 Minor Version 1 Protocol | 2010 | tsv | 42 |
| 1345 | Character Mnemonics and Character Sets | 1992 | app | 41 |
| 8446 | The Transport Layer Security (TLS) Protocol Version 1.3 | 2018 | sec | 40 |
| 5545 | Internet Calendaring and Scheduling Core Object Specification (iCalendar) | 2009 | app | 35 |
| 3261 | SIP: Session Initiation Protocol | 2002 | rai | 33 |
| 5905 | Network Time Protocol Version 4: Protocol and Algorithms Specification | 2010 | int | 32 |

Table 1: Top 10 RFCs by errata filing count

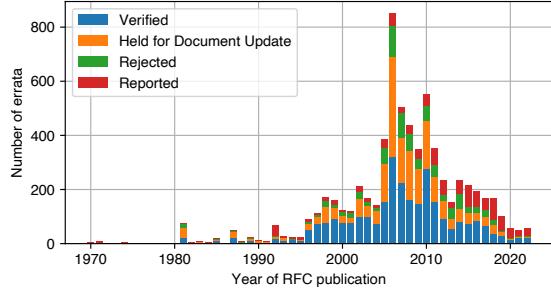


Figure 12: Errata filings by status, by publication year of the RFC.

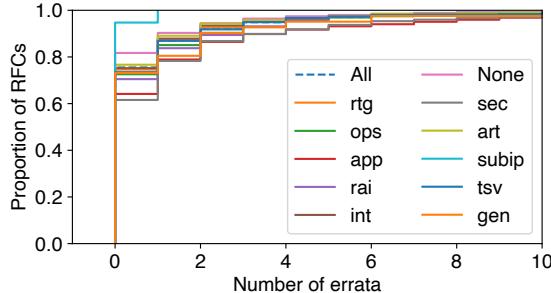


Figure 13: Cumulative distribution of errata filed per RFC by grouped by IETF area. The areas are as described in Figure 1.

update errata filed against them, only 40% have been updated or obsoleted by a subsequent RFC. We flag that this may be a cause for concern, or at least a missed opportunity for improvements to standards.

The third largest category (13%) is *rejected*, which covers errata that are invalid (like *erratum 6323*, which was rejected because the original text was understood to be correct) or proposes a significant change to the RFC that should be done by publishing a new RFC (like *erratum 5814*, which was rejected for proposing a significant change, rather than reporting an error). Such a large fraction of rejected submissions is unexpected and may flag issues with people’s understanding of the errata process and its place within the wider standardisation process.

Finally, 14.2% of errata are *reported* but unverified. Again, we are surprised to see unverified errata from over a decade ago, suggesting the process should be expedited.

Errata per RFC Area, Status, and Stream: Figure 13 shows a CDF of the number of errata filed, per RFC, in each IETF area. Non-IETF RFCs, e.g., IRTF and independent stream RFCs, and legacy IETF RFCs, are labelled as “None”. We confirm errata in standards are common: of the 4,373 standards-track RFCs in our dataset, 32.7% have attracted at least one erratum filing. However, there are three notable outliers:

1. RFCs published by the Sub-IP (*subip*) Area have very few errata, with only 5% of *subip* RFCs attracting errata filings. This temporary area – established in 2001 and concluded in 2005 – only published 19 RFCs, resulting in a far smaller sample than the other areas. For comparison, the next smallest area, General (*gen*), published 39 RFCs. *gen* RFCs attract a greater number of errata on average, vs. *subip* RFCs, likely due to their broader relevance.
2. We see that both the Application (*app*) and Security (*sec*) Area RFCs are more likely to have errata filed for them than other areas, with 35.9% of Application and 39% of Security RFCs attracting at least one erratum filing.
3. Table 2 details the split between *technical* and *editorial* errata across each area. While there is broadly an even split, there are areas where one type of errata is more dominant. For example, in the Routing (*rtg*) area, 60.8% of filings are editorial, while in the Applications (*app*) area, 59% were technical. It remains to determine why this is the case, and, in particular, to establish whether there is something inherent about the RFCs published by these areas that makes them more prone to containing errata, and to containing one type of errata vs. another. For example, in the Routing area, structured notation is frequently used to define routing entities; editorial errata are often filed in those definitions. Targeting such areas with improved alternate review procedures may be beneficial.

Tables 3 and 4 further categorise errata by the stream and status, at the time of publication, of each RFC. As expected, the majority of errata are filed against IETF RFCs and *Proposed Standards* since these make up the majority of RFCs that are published. However, there are notable differences in the average number of filings per RFC. *Proposed Standards* (1.01 errata per RFC), *Draft Standards* (1.93), and *Internet Standards* (2.17) attract a far higher number of errata per RFC than *Informational* (0.52) or *Experimental* (0.39) documents. This may be due to the additional

| Area | # | Verified | Held | Rejected | Reported | Technical | Editorial |
|---|------|----------|------|----------|----------|-----------|-----------|
| None | 1883 | 895 | 505 | 197 | 286 | 930 | 953 |
| Internet (int) | 650 | 281 | 223 | 98 | 48 | 342 | 308 |
| Operations and Management (ops) | 558 | 311 | 113 | 67 | 67 | 297 | 261 |
| Real-time Applications and Infrastructure (rai) | 457 | 143 | 213 | 48 | 53 | 255 | 202 |
| Security (sec) | 888 | 291 | 265 | 115 | 217 | 447 | 441 |
| Routing (rtg) | 831 | 305 | 378 | 140 | 8 | 326 | 505 |
| Applications (app) | 787 | 370 | 175 | 116 | 126 | 464 | 323 |
| Transport (tsv) | 459 | 188 | 142 | 75 | 54 | 258 | 201 |
| General (gen) | 41 | 22 | 4 | 5 | 10 | 8 | 33 |
| Applications and Real-Time (art) | 204 | 64 | 32 | 16 | 92 | 143 | 61 |
| Sub-IP (subip) | 1 | 1 | 0 | 0 | 0 | 1 | 0 |
| All | 6759 | 2871 | 2050 | 877 | 961 | 3471 | 3288 |

Table 2: Errata statistics by area.

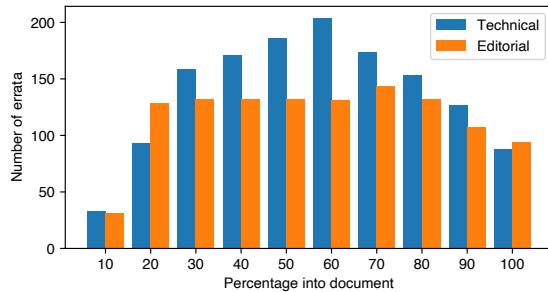


Figure 14: Errata counts by percentile location in document (0 is the beginning; 100 is the end).

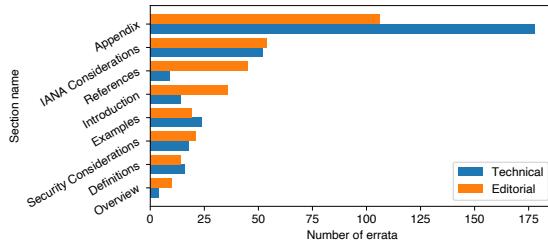


Figure 15: Errata counts by section title for the more frequent section titles.

readership and attention that standards-track documents receive, and because they are more likely to be the basis for future work and protocol extensions.

Errata Location: Finally, we investigate the location of errata within RFCs. Figure 14 presents the number of errata occurring at each decile of the documents, for the 2,552 filings where accurate location information is available, and after the copyright notice and other boilerplate has been removed. We see that technical errata dominate over editorial in almost all places, except for the very beginning where the *Introduction* is located. Moreover, it shows that the most technical errata are near the middle of the document where the most complex content is. We explore where errata occur, with Figure 15 showing section titles for errata appearing in at least 10 documents. Sections such as the *Introduction* or *References* are dominated by editorial errata while more technical sections,

like *IANA Considerations* (i.e., parameter registrations), *Security Considerations* or *Definitions*, have a larger proportion of technical errata. In addition, we see that sections labelled *Appendix* attract a significant proportion of technical errata. While appendices vary in their content, they are widely used to provide pseudocode and test vectors, or to describe algorithms. This suggests that it may be useful to target review efforts on appendices and other dense technical content where errata are more likely.

3.3 Summary

On RFC Production and Complexity (§3.1) we see evidence that Internet standards are maturing and that the pace of innovation has slowed. The rate of standards publication peaked in the mid-2000s, and has gradually declined since, and the number of IETF working groups has also started to decline. The standards that are being produced are taking longer, with more revisions prior to publication, and make increasing use of normative language and references to prior work; all evidence of complexity and the burden of backwards compatibility. The signs are that the IETF has developed, and is maintaining and extending, the core Internet standards, but has not broken out of that niche to new topic areas.

The RFC errata (§3.2) process seems to be broadly effective at finding problems, although there are a concerning number of unresolved errata reports that are perhaps indicative of an organisation that doesn't effectively maintain supposedly completed work. The high prevalence of errata in security-related RFCs also deserves further scrutiny to determine whether it is due to a higher number of defects in these RFCs or because defects are more likely to be noticed because these RFCs receive more review.

4 Trends in Participant Demographics

In this section, we look at the authorship of published RFCs to explore how participation in the IETF has changed over time. We show that the geographic distribution of RFC authors has shifted over time, with the proportion of authors from the US decreasing and a corresponding rise in authors from Europe and China (§4.1). We similarly explore shifts in affiliation, showing increasing influence of Huawei and ongoing shifts in the fortunes of many prominent US and European technology companies (§4.2).

| Stream | # | Verified | Held | Rejected | Reported | Technical | Editorial |
|-------------------|------|----------|------|----------|----------|-----------|-----------|
| IETF (6619) | 5797 | 2348 | 1815 | 798 | 836 | 3034 | 2763 |
| IAB (124) | 55 | 25 | 13 | 3 | 14 | 23 | 32 |
| Independent (376) | 330 | 235 | 32 | 28 | 35 | 172 | 158 |
| Legacy (1929) | 510 | 226 | 182 | 39 | 63 | 198 | 312 |
| IRTF (97) | 67 | 37 | 8 | 9 | 13 | 44 | 23 |
| All | 6759 | 2871 | 2050 | 877 | 961 | 3471 | 3288 |

Table 3: Errata statistics by stream; the “*Editorial*” stream has no documents, and is not shown.

| Status | # | Verified | Held | Rejected | Reported | Technical | Editorial |
|-----------------------------|------|----------|------|----------|----------|-----------|-----------|
| Proposed Standard (4084) | 4142 | 1680 | 1308 | 555 | 599 | 2213 | 1929 |
| Informational (2894) | 1500 | 719 | 399 | 136 | 246 | 754 | 746 |
| Internet Standard (147) | 319 | 118 | 111 | 66 | 24 | 135 | 184 |
| Best Current Practice (316) | 233 | 111 | 53 | 30 | 39 | 80 | 153 |
| Historic (70) | 20 | 13 | 4 | 2 | 1 | 9 | 11 |
| Draft Standard (142) | 274 | 94 | 93 | 66 | 21 | 150 | 124 |
| Experimental (563) | 221 | 115 | 65 | 20 | 21 | 121 | 100 |
| Unknown (929) | 50 | 21 | 17 | 2 | 10 | 9 | 41 |
| All | 6759 | 2871 | 2050 | 877 | 961 | 3471 | 3288 |

Table 4: Errata statistics by status at publication.

4.1 Geographic Distribution of Authors

The IETF Datatracker maintains metadata about document authors, including their names, email addresses, affiliations, and location information. This dataset does not cover the entire RFC corpus, with metadata available for authors of RFCs published from 2001, and where it has been provided country data is available for around 70% of authors, while affiliation information is provided for around 80% of authors.

Figure 16 shows how the proportion of RFC authors from different countries varies over time, with Figure 17 breaking this down by continent. The IETF has signalled that it wishes to encourage greater geographical diversity [26, 4]. Without an explicit goal, we frame our findings within the context of global population distribution. We find that North America, while still disproportionately over represented, is becoming less dominant. 75% of authors were from North America in 2001, and this has declined to 44% in 2020. At the same time, representation of both Europe and Asia has grown, from 17% to 40% and 6% to 14%, respectively. However, Africa and South America remain heavily under-represented, with only $\approx 0.5\%$ of authors coming from either continent in 2020. This suggests that, if the IETF is to become more geographically representative, further efforts are needed.

4.2 Author Affiliations

To explore trends in author affiliations, we gather affiliation data from the IETF Datatracker. This data is processed to normalise affiliation names by removing common variations in spelling and to amalgamate known subsidiaries and merged companies. For example, Huawei and Futurewei are combined as Huawei, and Sun Microsystems is merged with Oracle.

Corporate Authors: Figure 18 shows the top ten affiliations by proportion of RFC authors each year. We observe several

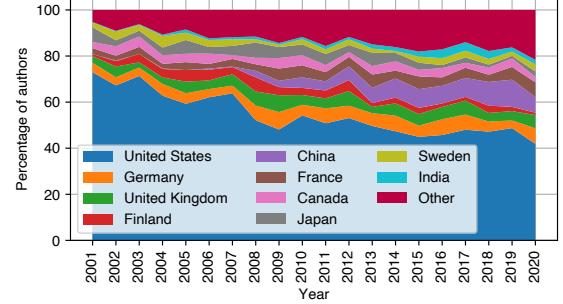


Figure 16: Authorship countries (normalised)

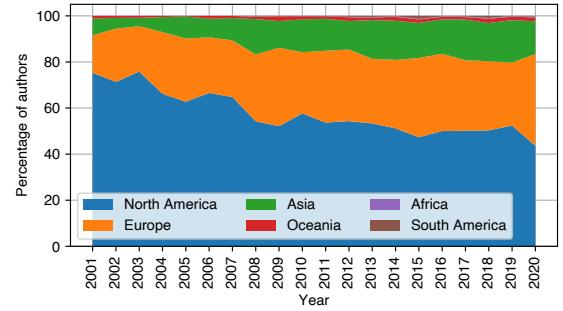


Figure 17: Authorship continents (normalised)

interesting trends. First, Cisco remains a consistent employer of IETF contributors, with around 12% of authors affiliated with the company in 2020, and having been the single largest affiliation across all years in the dataset. We can also see the rise of Huawei beginning in 2005, with 7.1% of authors being affiliated with the company in 2020, having peaked at 9.7% in 2018. Google has a similar trajectory, first appearing in the dataset in 2006, with 3.8%

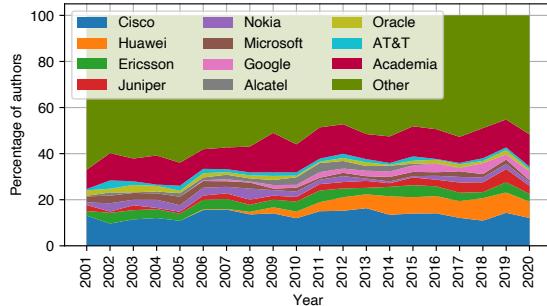


Figure 18: Authorship affiliations (normalised)

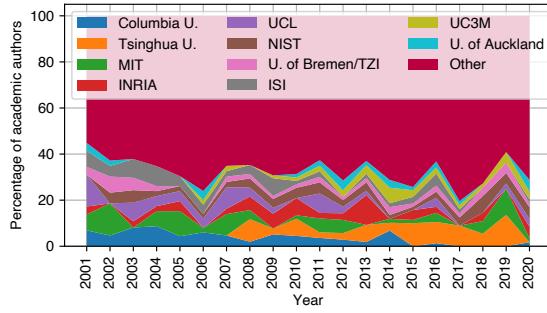


Figure 19: Academic affiliations (normalised)

of authors being affiliated with it in 2020.

We also observe the decline of a number of affiliations. Microsoft and Nokia, having peaked with 3.3% and 3.6% of authors, had 0.7% and 1.7% of authors in 2020, respectively, with the absolute number of authors from both companies also declining.

These shifts in affiliation reflect changes in fortune and commercial success of the companies. They also show the ongoing importance of IETF standards, with new companies choosing to send participants: while we do not demonstrate any causal link, it is encouraging that commercially successful companies opt to enable their employees to actively participate in the IETF. Care is needed to ensure that this relevance is maintained, however: the author pool has grown less diverse in terms of companies that are represented. 35.4% of authors came from the top 10 affiliations in the dataset in 2020, compared with 25.6% in 2001.

Academia and consultants: Academic affiliations are those where the affiliation name contains “University”, “Institute”, or “College”, while consultancy affiliations are those that contain “Consultant” (we recognise that this definition of consultants is limited and incomplete). Affiliation data has been normalised to remove common abbreviations (e.g., “U.” for “University”) and to translate non-English affiliations. As shown in Figure 18, we find that an increasing number of authors come from academic affiliations, growing from 8.1% of authors in 2001, to 13.6% in 2020, having peaked at 16.5% in 2009. The number of consultants, by our limited measure, has remained stable, accounting for 2% of authors in 2020. The remaining authors are largely from industrial affiliations.

Figure 19 shows the top 10 academic affiliations in the dataset, and the percentage of academic authors that have those affiliations

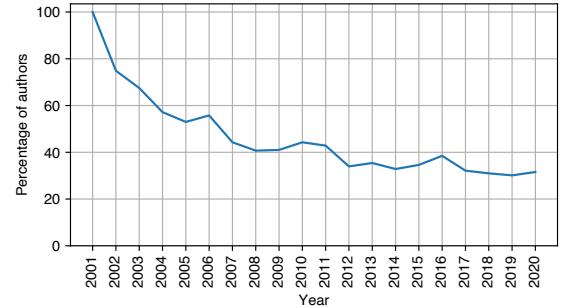


Figure 20: Percentage of new authors per year.

over time. In general, academic affiliations are each typically held by a small number of authors. We can see a number of trends in academic authorship, with fewer authors from Columbia University, MIT, and ISI in recent years, and the rise of Tsinghua University.

While some academic affiliations are shared by large groups of participants (e.g., MIT, ISI, Tsinghua University), academic participation is often driven by a single research group or individual within an institution. This allows smaller institutions, such as the University of Bremen and the University of Auckland, that appear in Figure 19, and others with single but prolific authors, to have outsize influence. Understanding whether this indicates academics pushing for standardisation of the results of government funded research to increase its impact, the outcomes of industry sponsored research, or academics acting as consultants for industry, is for further study.

Arrival of new authors: Figure 20 shows the percentage of authors in each year that have not previously authored an RFC. Given that the dataset used here begins in 2001, 100% of authors are new in that year. The more stable trend in recent years likely highlights the churn in RFC authorship, with around 30% of authors each year having never previously authored an RFC. This is consistent with the presence of authors with new affiliations and from new countries and shows an ongoing influx of new participants.

4.3 Summary

The trends highlighted in this section indicate a pool of authors that is slowly diversifying and changing over time, with a growing proportion of authors from outside of North America, contributions from new companies and affiliations, and relatively high authorship churn. Despite this, we find that certain regions and groups are not well represented in the IETF (e.g., contributors from Africa), and that the authorship pool is becoming increasingly centralised, with a third of authors coming from the top 10 affiliations. This suggests that, if the IETF is concerned with being more representative, further efforts are needed.

5 Organisational Dynamics of the IETF

While the IETF does hold regular plenary and interim meetings, much of the interaction between participants takes place on the public mailing lists of the various working groups. In the following

(§5.1), we characterise the mailing list interactions considering the volume of discussion, the number of drafts discussed, the duration of individual participation, and how interactions between participants change over time.

We then consider influence and impact of participants (§5.2). We measure how centralised is the active IETF community and how reliant is it on a small core of participants; how the most influential participants behave; how influence, determined by mailing list participation, relates to wider impact in the IETF; and whether the organisational affiliation of participants influences adoption of new work.

Finally, we consider the hierarchy and communication dynamics in the IETF (§5.3). We ask whether participation in IETF leadership roles is growing or becoming more centralised onto a small group of participants; whether the formal organisational hierarchy is reflected in the patterns of communication between participants; how people in different roles communicate and whether information flows up or down the hierarchy; and whether contact with people in leadership positions influences mobility within the IETF.

5.1 Characterising Participant Interactions

The operations of the IETF is largely underpinned by mailing list interactions, which are used to discuss and finalise the drafts that eventually become RFCs. Our data, which starts from 1995, confirms their vital role, with 1,153 mailing lists, containing 2,439,240 emails from 74,646 unique email addresses.

Volume of Discussion: Figure 21 shows the number of emails sent across the last 25 years (dashed red line), showing that email volumes have grown significantly with time, plateauing at around 130,000 messages per year since around 2010. The figure also shows (solid blue line) how the number of unique participants (with entity resolution performed as discussed in Section 2.2) observed every year in the mail lists, showing a decreasing trend in the number of contributors since 2007. We note that the number of participants is broadly correlated with the rate of RFC publication (Figure 1), likely indicating that many participants disengage once the document they were working on is published.

Figure 22 breaks down the messages sent into four categories based on the sending address:

1. *Datatracker-mapped* addresses represent those participants sending email in a given year that have an IETF Datatracker account associated with their email address. Datatracker accounts are needed to submit drafts and register for meetings, and to allow those in leadership roles to perform administrative actions, so this indicates the fraction of participants who are strongly engaged with the IETF.
2. *Automated* addresses are those sent by automated systems rather than directly by participants. This includes the automated announcement of new draft submissions and meeting, last calls for comments, etc., sent by the Datatracker, but there is also a growing number messages sent by GitHub and other version control systems used for managing drafts. There were 122 active IETF working groups at the time of writing. Of these, 17 listed a GitHub repository in their metadata.

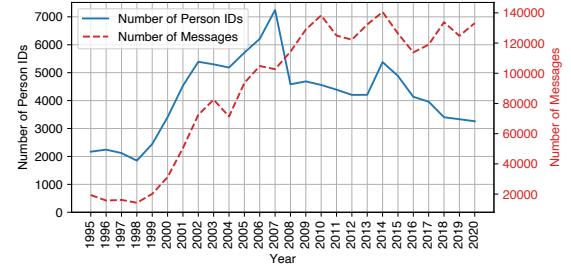


Figure 21: Number of email messages per year and number of unique participants exchanging emails per year.

The QUIC working group, as one example, has replaced the typical email list discussions with GitHub issues: indeed, this is a significant part of the surge observed in 2016. Given that most working groups use mailing lists to manage their activity, we do not further analyse interactions that take place on GitHub.

3. *Role-based* addresses represent those sent by people in formal leadership or administrative roles in the context of their role. There are a very limited number of these addresses, e.g., the IETF, IRTF, and IAB Chairs, the IETF Executive Director. As seen, these addresses are used sparingly. Note that working group chairs and area directors use their individual email addresses when managing their working group or area, and so appear as Datatracker mapped addresses.
4. *Unmapped* addresses are those that do not fall into the other categories. They typically represent participants that lurk on a mailing list and provide occasional input, but who are not document authors and do not attend meetings. This category has reduced over time as more types of interaction require a Datatracker account and as comments from occasional contributors are increasingly submitted via issue trackers on GitHub and similar services.

The increasing use of GitHub suggests that the data in Figure 21 likely understates the volume of interactions. The plateau observed in the number of messages sent in recent years is at least somewhat attributable to the shift to GitHub and similar services. As this shift continues, it will become important for future work to consider these interactions.

Discussion of draft documents: The mailing lists are, other than plenary or interim meetings, the primary means for the discussion of draft documents. To measure how often drafts are discussed, we identify mentions of Internet-Drafts and RFCs in mailing list messages. We extract any mention of a draft (beginning `draft-`) or RFC (i.e., “RFC” followed by a number). Figure 23 presents the number of drafts mentions in the emails per year. Separate mentions of the same draft are counted as different mentions, as we want to observe the entire volume of mentions. We observe a strong increase in the number of mentions over time.

When compared with the number of active drafts (Figure 3 in §3.1) we see that the number of mentions of work-in-progress drafts and RFCs is increasing faster than the number of such

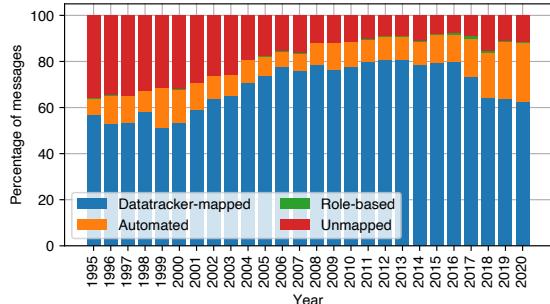


Figure 22: Number of messages exchanged per year across different categories: messages mapped to the Datatracker (Datatracker Person-ID); messages by automated email addresses (Automated); messages by role-based addresses (Role-based); and messages not mapped to the Datatracker (New Person-ID)

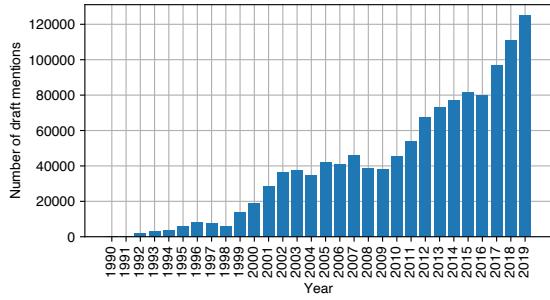


Figure 23: Number of mentions of Internet-Drafts and RFC in the IETF email archive, per year.

documents. This further supports the hypothesis of increasing standards complexity—discussion is increasingly referencing prior standards and other work in progress, as would be expected if the task of ensuring compatibility with that work is growing in complexity.

Duration of Participation: We now look into the duration of participation, in terms of the number of years that people actively participate in any IETF mailing lists.

We define the *contribution duration* of a participant as the length of time that they have contributed to the mailing lists. To do this, we study the participants who first send email to mailing lists between the years 2000 to 2013. We limit our analysis to 2013, since the longevity of contributors that first participate more recently cannot be determined. For each year, in the period 2000 to 2013, we look at those who first sent email in each year and the number of years that they then go on to remain active in the mailing lists. For example, a participant who first sent an e-mail to an IETF list in 2010, and last sent an e-mail to an IETF list in 2018, will have a contribution duration of 9 years.

First, to understand duration of participation, we first generate Gaussian Mixture Models for observing different clusters of the maximum duration of contribution. These models identified three broad clusters into which the participants can be categorised: *young contributors* for whom the time between their first and last emails

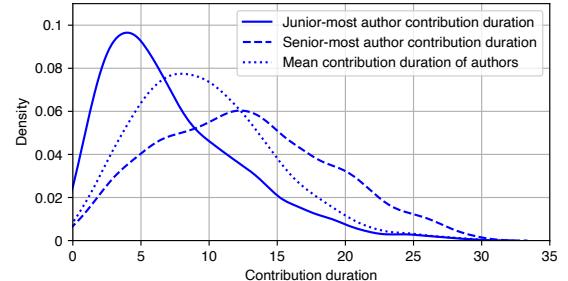


Figure 24: Contribution duration distribution of authors of RFCs: junior-most author of each RFC, senior-most author of each RFC, and mean contribution-duration of all authors of each RFC

to the IETF lists in one year or less; *mid-age contributors* who go onto remain active participants in the email lists for more than one year, but less than five years; and *senior contributors* who remain participants in the mailing list for five or more years.

Next, to begin to understand how contribution duration relates to RFC authorship, we look at the email interactions of the authors of each RFC. Specifically, we look at the distribution of the contribution duration of the authors of RFCs, considering the *mean contribution duration* of all of the authors of the RFC at the time of publication, the contribution of the *junior-most author* (the author with the lowest contribution duration at the time of publication), and the contribution of the *senior-most author* (the author with the highest contribution duration at time of publication). Figure 24 shows the distribution of contribution duration of each of these three measures. This shows that the majority of junior-most authors have participated for less than 5 years in the IETF (i.e., they are young or mid-age contributors), whereas the majority of senior-most authors are senior contributors who have participated in the IETF for significantly more than 5 years (in fact, 35% of authors exceed 15 years of IETF participation). This shows that RFCs tend to be authored by a mix of seniority levels.

Evolution of interactions: The analysis in Section 3.1 showed that RFCs are taking longer to publish and show evidence of increasing complexity. Such increased complexity should be visible in the email discussion during preparation of the RFCs.

Figure 25 explores this, showing the change in the number of other people RFC authors interact with (i.e., the annual degree of their communication graph) over time. The amount of interactions RFC authors engage in has substantially increased over time. For instance, in the year 2000 only 5.5% of the authors had a degree of over 25 (i.e., they exchanged with more than 25 other people), whereas by the year 2015 almost a quarter of the authors had a degree over 25. This confirms that, on average, more recent RFCs generate greater discussion. This helps to explain increasing publication times: as documents become more complex, authors spend more time interacting with other participants to resolve the complexity, and these discussions are likely to lead to more drafts.

We also contrast the interaction patterns of the junior vs. senior authors. Figure 26 presents the CDFs of the in-degree (i.e., number of messages received) across both junior and senior authors. It shows that the incoming interactions from senior contributors to

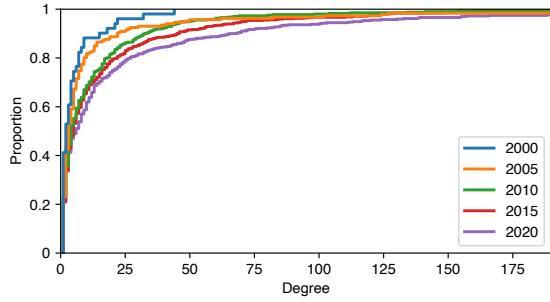


Figure 25: CDF showing drift in annual degree (interaction with their network) of RFC authors over the period 2000-2020.

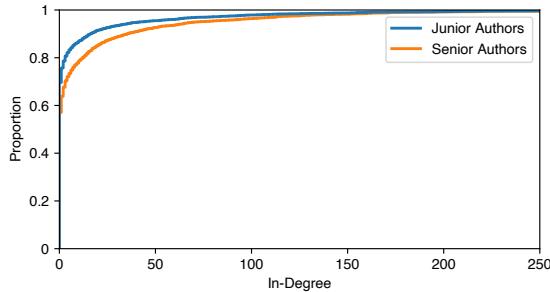


Figure 26: CDF showing number of senior contributors sending messages (in-degree) to junior and senior authors.

junior authors are significantly less than the incoming interactions from senior contributors to senior authors. Nearly 55% of junior authors receive messages from fewer than 10 senior contributors, whereas nearly 65% of senior authors receive messages from more than 10 senior contributors. This speaks to the differing roles played by these sub-populations: senior authors act as hubs through which substantial volumes of interaction flow.

Summary: Emails play a vital role in underpinning RFC publication with approximately 130,000 emails sent per year. We have shown that the seniority of both participants and RFC authors fundamentally changes the volume of interactions that they have. These trends are likely to have implications for the IETF community, especially as it tries to encourage new participants.

5.2 Characterising Influence and Impact

As we showed in Section 5.1, protocol standardisation is an inherently social process with most day-to-day work happening on public mailing lists. We are specifically interested in better understanding how influence is distributed across stakeholders and how it might affect the standardisation process. This is of critical societal importance: the IETF has a major impact on global Internet technologies, and understanding the social processes involved would give us insight into not only the driving forces behind standardisation, but also its resilience to the loss of major participants. In this section we consider the following:

1. How centralised is the active IETF community, and to what extent is it reliant on a small core of participants?

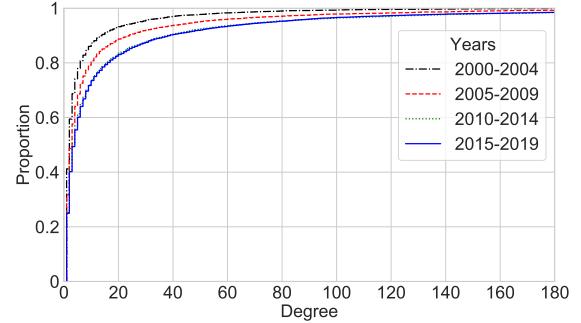


Figure 27: Cumulative degree distribution of the email graph for different year-periods.

2. How do the most influential participants behave?
3. How does influence (determined by mailing list participation) relate to wider impacts throughout the IETF?
4. Does the organisational affiliation of participants also influence the innovation (adoption of new work) within IETF?

5.2.1 Measuring Influence: The IETF as a Social Graph

To measure influence of participants in the IETF we extend the analysis in Section 5.1 and, for each year, build a social graph based on the *active community* of participants who have email interactions with any other participant in the previous five years.

We previously demonstrated that there are three categories of participants in IETF: *young contributors* who leave within one year of their first year mailing list contribution; *mid-age contributors* who stay active for up-to five years; and *senior contributors* who remain active for more than five years (§5.1). Following this, a 5-year period window is chosen to observe interactions. There are no direct participant-to-participant email exchanges in our dataset: the emails and responses we capture are sent to the public mailing lists, and we observe an interaction between two participants when one replies to an email sent by another on any mailing list. This yields a social graph based on 1,049,793 emails (out of 2.1M; not all messages receive a reply) from 22,138 unique participants across 840 mailing lists.

Participation To examine the variation in participation between IETF participants, in Figure 27 we plot the cumulative degree distribution of the social graph (i.e., the cumulative number of people with which a participant interacts). We observe a core group that always interacts with substantially more people than the rest of the community, with around 80% of the participants interacting with less than 15 others (i.e., having degree < 15), but 5%-10% interacting with more than 40 others. This difference decreases over time: in 2005-2009 only around 6% of participants had degree >40, but by the period 2015-2019 this increases to around 10%. While there is always a core of more active participants, the amount of interaction has increased over time. This confirms the results in Figure 25, for RFC authors, and shows that they also apply to interactions in the community at large.

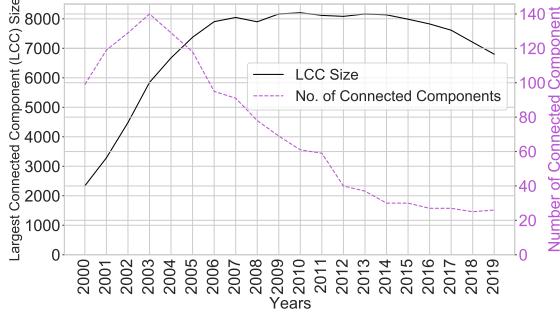


Figure 28: Size of Largest Connected Component (LCC) and Number of Connected Components (NCC) of the email graph.

To understand the structure of the community, and its reliance on specific groups, we analyse the connected components of this graph. Each connected component reflects a maximal set of nodes such that each pair of nodes is connected by a path. We compute the size of the Largest Connected Component (LCC) and the Number of Connected Components (NCC) for each year in Figure 28. The NCC peaks in 2003 before declining, and broadly aligns with the variation in the number of meeting attendees,⁸ suggesting a lag between participating for the first time and integrating into the wider community. In contrast, the size of the LCC increases until 2006 before stabilising. Overall, we observe that the IETF community has become less fragmented over time.

Influence: Relying on a small group to interconnect the community could undermine the resilience of the IETF. To study the influence of participants and their role in interconnecting the wider community, we compute the betweenness centrality of each participant in a given time period [56, 98, 87]⁹ to determine the most influential participants in the email graph. Figure 29 then shows the effect of removing the most influential participants (in 1% increments from most to least, moving left-to-right on the x-axis) on the size of the LCC. Worryingly, we find that removal of just 20-25% of the most influential participants causes the LCC to shrink by 90%. However, we also find that this impact has decreased over time: for instance, in 2000-2004, removing the top ~5% influential participants, reduces the size of LCC by more than half, whereas in 2015-2019, it takes the removal of the top ~15% of participants to have the same effect. This shows that the community has become *more* cohesive and resilient over time, and the IETF can now sustain a larger amount of churn while maintaining a well-connected social graph.

⁸<https://datatracker.ietf.org/stats/meeting/overview/>

⁹We considered other graph based influence metrics such as eigenvector centrality: eigenvector centrality reflects the importance of a node as per its neighbours, while betweenness centrality is based on shortest paths which is independent of the influence of neighbours. However, we found a very strong correlation between the two measures, in line with similar experiments [95, 42]: the Spearman's rank correlation of participants ranked by betweenness centrality and eigenvector centrality ranged between 0.51-0.72, with a strong statistical significance ($p < 0.01$) for the period 2000-2019. We therefore use just betweenness centrality in our analysis; this is often studied and acknowledged as a measure of influence in social and complex networks, particularly built on online communication [40, 24, 36].

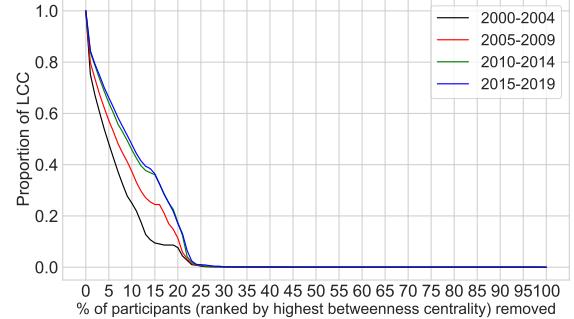


Figure 29: Impact of removing participants by their influence on the size of the LCC.

5.2.2 Behaviour of Influential Participants

We now characterise the behaviour of the most influential participants, as measured by betweenness centrality (§5.2.1), in terms of the volume of emails they send, the length of time they are active within the community, and the topics discussed.

Email volume: Figure 30 shows that each participant in the top 10% most influential participants sends on average around 0.05%-0.08% of total emails in a given period. Collectively, the top 10% most influential participants account for 43.75%, on average, of the total number of emails sent, a substantially larger proportion than the other participants.

This ratio seems broadly stable over yet the overall number of emails sent increases up to 2010 and remains roughly constant from then onward. Along with the results from §5.2.1, this shows a worrying, if slowly improving, dependence of the IETF on a small number of highly influential participants.

Cross-area review: In addition to sending more emails, we test if influential participants engage with different parts of the community more often than do typical participants. That is, we measure if influential participants are more likely to make contributions to working groups in multiple IETF Areas.

Figure 31 shows the mean number of areas where participants in the IETF are active (based on the set of mailing lists they send to). We see that more influential participants engage in more areas of the IETF, on average, than do less influential participants: not only do influential participants send more email, they send email to working groups in more areas of the IETF. This indicates that influential participants benefit the IETF in enabling cross-area discussion and review, and can bridge administrative divisions.

It can also be seen that cross-area engagement has improved over time, with participants as a whole sending email to groups in more areas over time. This shows that the IETF community as a whole is discussing more broadly than it used to, perhaps further indication of the increasing complexity of the standards development and widespread concern about feature interactions.

Participation duration: We next consider how long participants remain associated with the IETF, measured as the time between sending their first and last emails to one of the mailing lists.

Figure 32 shows the mean participation duration distribution for participants, ranked by influence. Participation duration increases

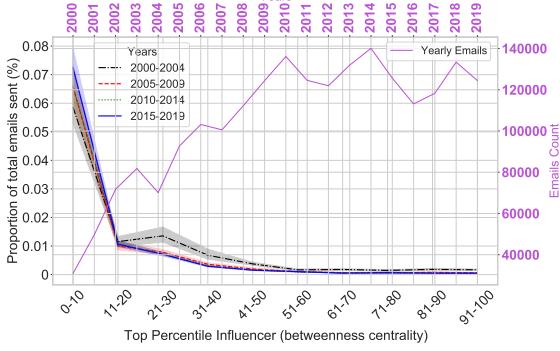


Figure 30: Proportion of emails sent (%) by participants in each period (with 95% confidence interval) according to their betweenness-centrality percentile (x-axis). Y2-axis (purple) shows yearly count of emails.

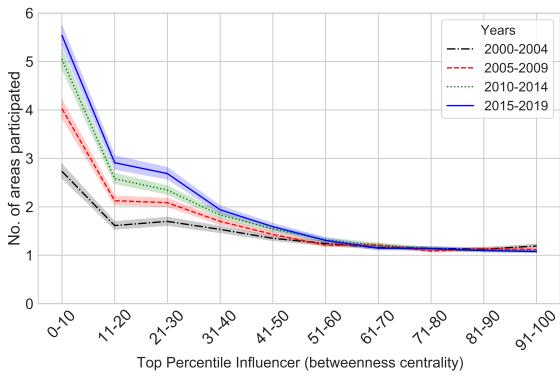


Figure 31: Mean number of areas participated in (with 95% confidence interval) according to their betweenness-centrality percentile in the x-axis, ranked from top (0-10) to bottom percentile (91-100).

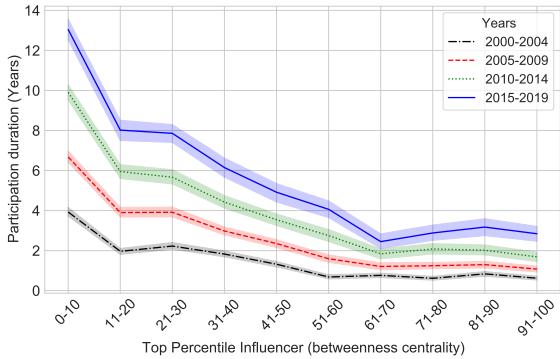


Figure 32: Mean participation duration (with 95% confidence interval) according to their betweenness-centrality percentile in the x-axis, ranked from top (0-10) to bottom percentile (91-100).

over time, with the most influential participants typically being those who have been active for longer. While this might be expected, it might be a good sign for the IETF community in that it implies that the most influential participants are also the most experienced.

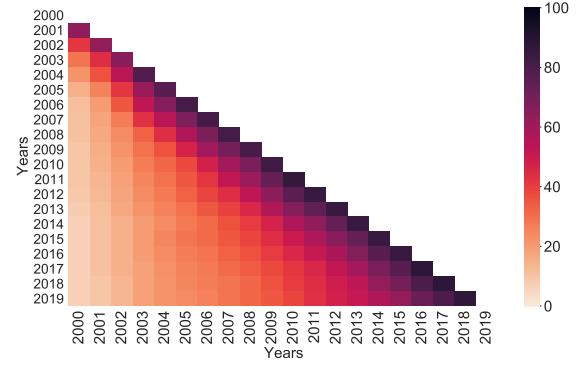


Figure 33: Yearly overlap (%) of top 10% influencers.

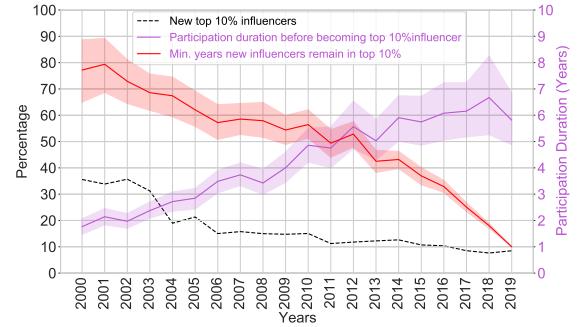


Figure 34: Percentage of new influential participants breaking into top 10% influencers (black); the average years of participation before entering top 10% (purple); and minimum # years new influencers remain in top 10% (with 95% confidence intervals).

The short participation duration of the less influential participants matches the participant churn discussed in Figure 20 in §4.2.

We further consider whether participants who become influential, remain influential. Figure 33 shows what proportion of the top 10% influential participants in any given year (y-axis) was also among the top 10% influential participants in any other year (x-axis). This shows that a significant majority of participants that become influential, continue to be influential for a number of years: the top 10% are influential for at least 6-7 years on average.

Finally, Figure 34 shows that breaking into the top 10% of influencers requires an increasing number of years of participation. This may be beneficial, showing that the IETF is maturing and is capable of retaining influential and experienced participants, but it may also point to an increasingly ossified structure that is not welcoming to newcomers.

Topics of Discussion: The topics discussed on working group mailing lists are a good indicator of the focus of technical work, and it might be expected that influential participants will set the direction of that work. To explore this, we use the Latent Dirichlet Allocation (LDA) [13, 45] model from gensim [80] to induce 100 topics on the entire set of email texts.

Each topic is a distribution over words. For example, a *security* topic might have high probability for *cipher*, *rsa*, *auth*, and related words. We can also use the model to obtain a vector for any input

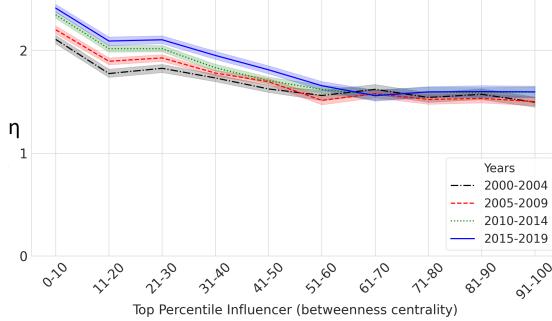


Figure 35: Mean topic entropy (with 95% confidence interval) according to their betweenness-centrality percentile in the x-axis, ranked from top (0-10) to bottom percentile (91-100).

text as a sparse distribution over all topics, e.g., finding that a text is comprised of 30% security and 70% video streaming topics. With this, we generate a vector for each participant in a given time period by concatenating all messages sent by that participant in the period and feeding it into the LDA model. The result can be interpreted as a distribution of topics for each participant.

Manual inspection of the induced topics reveals some interesting trends. We find that topics related to *routing* and *email* protocols are seeing a steady decline in popularity, while topics like *streaming* and *cryptography* are becoming more prominent. This reflects wider trends in standardisation, as these efforts have adjusted to the public’s increased awareness of privacy, especially in light of the Edward Snowden leaks [32, 31]. We also find that as videoconferencing became increasingly popular, so did the WebRTC protocol that frequently underpins them.

We next measure diversity of a participant via the entropy of their topic vector, their *topic entropy*, defined as follows:

$$\eta = - \sum_{i=1}^{|T|} p_i \log p_i$$

where $T = [p_1, \dots, p_N]$ is a topic vector, which defines a probability distribution over N topics, each p_i is a probability of the i -th topic and $\sum_i p_i = 1$.

Figure 35 shows topic entropy distributions of participants in different time periods and influence percentiles. While we initially experimented with measuring diversity by simply counting the number of topics that account for the majority ($\geq 95\%$) of a participant’s topic distribution probability mass, we found that most participants engage in a relatively small number of topics, regardless of influence. However, after measuring diversity using entropy, we identify that the activity of the more influential participants tends to be more evenly spread across the topics they participate in¹⁰. This difference becomes more pronounced over time, implying that increased topic entropy (i.e., more evenly participating in different topics) is an increasingly salient property of influencers. This is aligned with our earlier findings on cross-area review, showing the growing number of areas in which influencers participate.

¹⁰An example with 3 topics: a distribution of [0.8, 0.1, 0.1] is less evenly spread than [0.3, 0.4, 0.3].

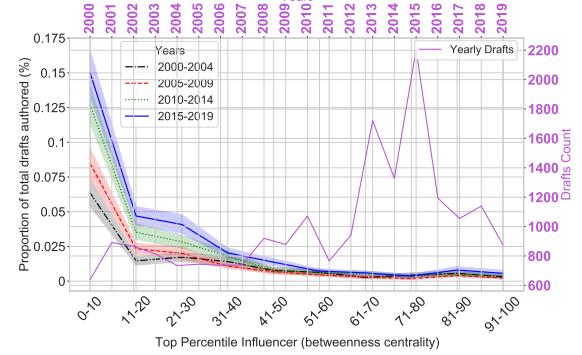


Figure 36: Average drafts authored (with 95% confidence interval) according to their betweenness-centrality percentile in the x-axis, ranked from top (0-10) to bottom percentile (91-100). Right-y axis (purple) shows yearly drafts per year.

5.2.3 Impact of Influential Participants

While the top participants are defined by the betweenness centrality metric (§5.2.1) as being influential within the IETF mailing lists, it is unclear how this influence translates into document authorship and leadership roles within the IETF. We explore these aspects of influence in the following.

Document authorship: The output of the IETF is technical standards and other documents published in the RFC series. RFCs are developed by working groups from a sequence of Internet-drafts, with individuals acting as named authors. Having identified the most influential participants in the mailing list community, we can determine whether or not these same participants are also the most active authors.

Figure 36 shows the distribution of the proportion of documents authored by mailing list participants, sorted by their influence rank. We observe that: (i) influential participants in the email graph, i.e., those ranked top in betweenness centrality, tend to write more drafts than do participants with lower centrality; and (ii) the number of drafts per participant is increasing over time for all participants. For example, during 2010-2019, each participant in the top 10 percentile, authored 0.125% to 0.175% of the total number of drafts in that period. We find that 32%-42% of the total drafts, in different years during this period, were authored by the top 10% most influential mailing list participants.

Document co-authorship and correlation with the email graph: Figure 36 showed that influential participants on the mailing lists tend to author more drafts than others. We next ask if they are also influencers in the draft co-authorship graph.

To do this, we create a draft co-authorship graph where each author is a node, and draft co-authorship is an edge. We then measure influence with betweenness centrality of the authors in each time period. Table 5 compares the top 20% of influencers of the mailing lists with those of the co-authorship graph. We find a significant overlap between both groups, ranging from 48.2% to 67.2%. There is a significant ($p < 0.05$) positive correlation between the rankings of the overlapping members of each community: participants that are influential in the mailing lists are also

| Top 20% draft authors & all email participants | | | | | |
|--|------------------|-------|---------|------------------|----------|
| Years | Sub-Network Size | | Overlap | Spearman's r_s | |
| | Co-author | Email | | r_s | p-value |
| 2000-04 | 398 | 1390 | 48.24% | .323 | 4.75e-06 |
| 2005-09 | 427 | 1662 | 67.21% | .332 | 7.39e-09 |
| 2010-14 | 728 | 1639 | 63.05% | .299 | 5.73e-11 |
| 2015-19 | 915 | 1370 | 55.85% | .337 | 4.13e-15 |

Table 5: Overlap in the co-author and email graph.

likely to be influential in draft authorship.

We also look at the top 20% participants from the co-authorship network who are *not* part of the top 20% influencers in the email network in Figure 37. We observe that some prolific authors are *not* engaged in the email discussion: 40%-50% of non-overlapping authors are ranked between 20th to 40th percentile of influence in the email network. These non-overlapping authors are typically more junior with respect to their participation duration.

Influencers and leadership roles: Working group chairs are selected from the community, so we might expect those selected to be influential in the community. Figure 38 shows that this is indeed the case: 87.5% of working group chairs are in the top 20% of mailing list influencers and 67.4% are in the top 20% of document authors in the year before they became chairs. This also shows the impact of taking up a leadership role: influence in both the mailing list and authorship communities grows in the year after participants become a working group chair for the first time.

Influence of organisations: While participants in the IETF are expected to contribute as individuals, they are usually affiliated with an organisation. To study the potential influence of those organisations, for each draft we obtain the authors' affiliations from the Datatracker. If this is not available, we use the domain name of the authors' email addresses, which we map to the relevant organisation (e.g., @cisco.com maps to Cisco). For generic email addresses, such as @gmail.com, we use the participant name if no affiliation is available.

Figure 39 shows the ten most frequent affiliated organisations of the top 20% of draft authors in the period 2015-2019, and the number of authors affiliated with these organisations in each period. While the dominance of Cisco is clear, other organisations, such as Huawei, have gained a larger presence over time. In general, a small number of organisations employ a significant fraction of the influential participants in the standards process. If we look at the period 2015-2019, for example, out of the 915 authors ranked in the top 20% of most influential authors, 342 belong to the ten most frequently affiliated organisations, and nearly 253 are from Cisco, Huawei, Ericsson, or Juniper alone.

While authors from the same organisation often co-author drafts together, collaborations between authors from different organisations are also common. Figure 40 shows the most common collaborations between authors from the top 15 most frequently affiliated organisations. Collaborations between authors at competing organisations, such as between authors from Cisco, Huawei, Ericsson, and Juniper, occur frequently. We also computed that over the years, joint collaboration (authors from multiple affiliations)

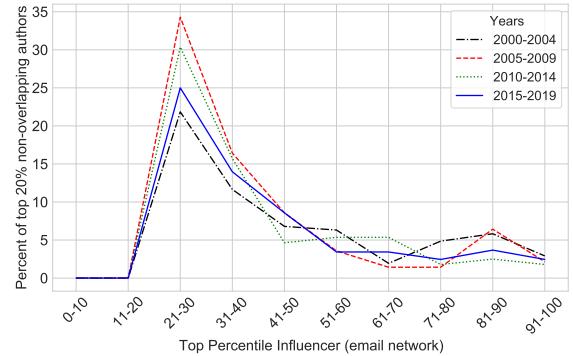


Figure 37: Percentage of Top 20% authors that are not 20% influencers according to their betweenness-centrality (email graph) percentile in the x-axis, ranked from top (0-10) to bottom percentile (91-100).

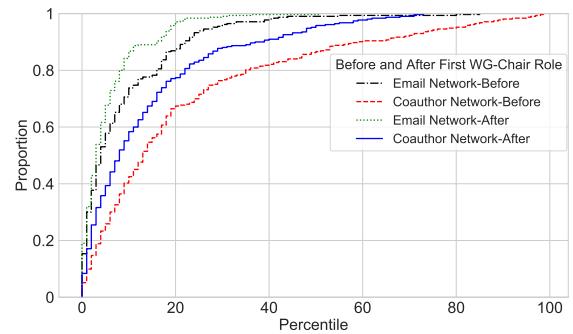


Figure 38: CDF of the percentile of betweenness-centrality of WG chairs in the email and co-authorship graph, one year before and after becoming chairs.

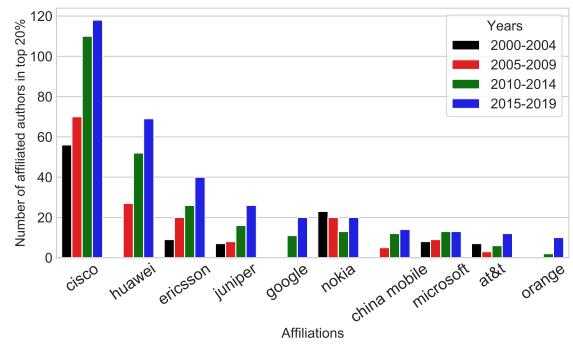


Figure 39: Authors in the most common affiliations of top 20% authors (co-authorship graph) in 2015-2019.

has increased: 772 co-authored drafts in the period 2000-2004 vs. 3083 co-authored drafts in the period 2015-2019.

Further inspection of the collaboration trends reveals that the volume of collaborations between authors from Huawei and other organisations appears to be unaffected by the addition of Huawei to the U.S. Entity List [94] as the trends in Figure 40 also hold for 2020-2021.

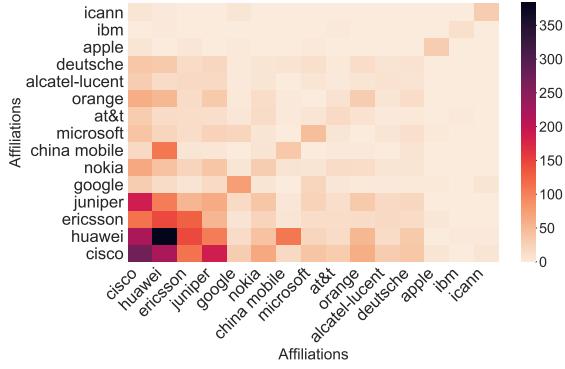


Figure 40: Number of drafts co-authored by top 15 most frequently affiliated organisations. We only include the top 20% authors (co-authorship graph) of 2015-2019.

5.2.4 Summary

The IETF is a reasonably centralised organisation, that dominated by a relatively small group of influential participants, but this dominance is gradually reducing over time. However, removing around 20%-25% most influential participants still fragments the email graph (§5.2.1, Figure 29).

The most influential participants are those with the highest degree of engagement within the community. These also tend to be more involved in draft authoring (§5.2.2, Figure 36). However, we also observe that over the years more people, even those less influential in the email network, are getting involved in draft authoring activities (Figure 36). This shows that these activities, while centralised to a certain degree, are still open to contributions from the wider IETF community.

We observe that the influential participants show a much higher level of engagement with the community (§5.2.2). A considerable proportion of total emails is sent by the top 10% of influential participants (Figure 30). Compared to the rest of participants, they are active in more areas (Figure 31), are active within the IETF for a longer period of time (Figure 32), and participate in a more diverse set of topics (Figure 35). Finally, their influence extends from the mailing lists to other activities such as draft authorship.

A significant overlap and correlation is observed between influential authors from co-authorship network and email networks (§5.2.3, Table 5). This shows that a large set of authors exhibit an ability to co-author drafts as well as engage with participants on the email networks, thereby translating their influence from email networks to co-authorship and vice versa.

Finally, a considerable portion of influential participants (~30%) are affiliated with one of the more prominent organisations, e.g., Cisco or Ericsson, (Figure 39). Participants from different organisations do considerably collaborate (Figure 40). Moreover, the level of such collaboration is found to have increased with time.

5.3 Hierarchy and Communication Patterns

As a voluntary and consensus-driven [81] standards development organisation, one might expect the communication patterns in the IETF to differ from those in businesses following a more traditional

management structure. Those in IETF leadership roles, whether working group chairs or area directors, must lead by persuasion and example and have little in the way of coercive power compared to a typical manager and this can be expected to impact how they communicate.

In the following, we quantitatively determine the effect that organisational hierarchy levels have on communication patterns throughout the IETF. We study how the email communication is impacted by the hierarchical roles of individuals. According to the IETF mission statement [1], the IETF is made up of volunteers who collaborate to develop consensus-based technical standards. This collaboration should be evident in their communication practises.

We consider the following questions:

1. Are higher levels of the hierarchy growing or becoming more centralised? Are they associated with an increased domination of the conversation? (§5.3.1)
2. What is the association between the organisation hierarchy and general communication patterns? (§5.3.2)
3. How do people with differing roles communicate with each other? Does information flow up or down the hierarchy? (§5.3.3)
4. What is the impact that individuals have on their direct contacts' activity over time? How does this vary across hierarchy level? (§5.3.4)

Question 1 relates to the *steepness* of the hierarchy [2], that may exhibit a high amount of centralised decision making, compared to a more *diffused* [47] hierarchy where a larger amount of people in higher levels stifle lower level voice by creating “bystanders” who find it hard to speak up. In general, the conclusions about steep hierarchies are confused, where some papers conclude a positive and some a negative effect of steep hierarchies. We explore what this means for the IETF in §5.3.1.

Questions 2-4 relate to the *power distance* [60, 30, 39] of individuals, which is the number of hierarchy levels between people who are communicating. The focus is on the impact this has on people's voice or lack thereof when the distance is large versus small. The conclusions from this research area usually suggest that even small power distances have a large effect of reducing the voice of the lower of the two levels. We tackle this for the IETF in three different ways in §5.3.2 – §5.3.4.

The voluntary nature of the IETF offers a new perspective on these questions. We demonstrate that despite the hierarchy becoming more diffused, participants in the higher levels in the IETF hierarchy perform a facilitator role, promoting discussion with those they interact with and receiving a benefit themselves as a result. We also show that higher level participants tend to be a focus of discussion with remarks directed upward toward them whereas Regular Participants engage more in group-discussion with less focus on an individual. This suggests that the IETF does not conform to the notion that power distance has a negative effect on communication, nor that a more diffused hierarchy reduces the voice of lower levels.

5.3.1 Centralisation and Dominance

In this section, we look at the number of individuals that inhabit each level of the IETF hierarchy and their communication patterns, and quantify the evolution of the mailing list activity within each level of the IETF. We find that the middle level, Working Group Chairs, is gaining in overall proportion of activity, whilst activity of participants at the lowest level is decreasing. The top level of the hierarchy, area directors, remains constant in proportion. The increase in working group chair activity coincides with an increase in the number of chairs per working group.

Role of Working Group Chairs: Figure 41(a) shows the number of participants at each level of the hierarchy in IETF: Regular participants are the largest contingent, comprising 90–86% of active participants; this proportion is decreasing over time. The next level is Working Group Chairs, comprising 6–10% of active participants; this proportion is rising and becoming more diffused over time, with about 15 regular participants for every working group chair in the organisation in 2013 versus 8 participants per chair in 2021. The top level is Area Directors who consist of less than 1% of the total active population of the IETF; this stays constant over time.

The fraction of emails sent by participants at each level of the hierarchy is shown in Figure 41(b). Individuals at higher levels contribute a disproportionate share to communication and this share is increasing over time, consistent with §5.2. A smaller proportion of email are sent by regular participants at the same time as the number of working group chairs is increasing, further suggesting that the IETF hierarchy is becoming more diffused, in terms of mailing list communication, which may make cooperation between hierarchy levels more difficult.

Number of Working Group Chairs: Figure 42(a) shows the number of working group chairs over time. There is a 35% rise in the number of chairs while the amount of individuals that fill the chair roles remains mostly constant.

Figure 42(b) contains two estimates of the number of working groups over time, based on tracking group events (i.e., Datatracker metadata) and mailing list activity since groups may be recorded as being active but have no activity, and vice-versa. The discrepancy between the two estimates is the least (at most 10 working groups) from late 2014 to early 2019.

Figure 42(c) shows how many chairs exist per working group, using both estimates of the number of working groups from Figure 42(b). It is clear there is a slow rise in the number of chairs per working group over time. The ratio rises from between 1.7-1.9 to 2.0-2.2 within the middle 50%, late 2014 to early 2019. This perhaps represents a growing understanding in the IETF that it is desirable to have two or more chairs per working group in case of illness/unavailability or to avoid conflicts of interest. Therefore, this plot again shows that the hierarchy is becoming more diffused, even if the raw amount of individuals in the chair roles has remained constant. The multiple roles per individual may result in even more difficulty of regular participants to use their voice.

Number of Area Directors: A version of Figure 42 considering area directors rather than working group chairs is available, but omitted to save space as all three plots stay mostly constant

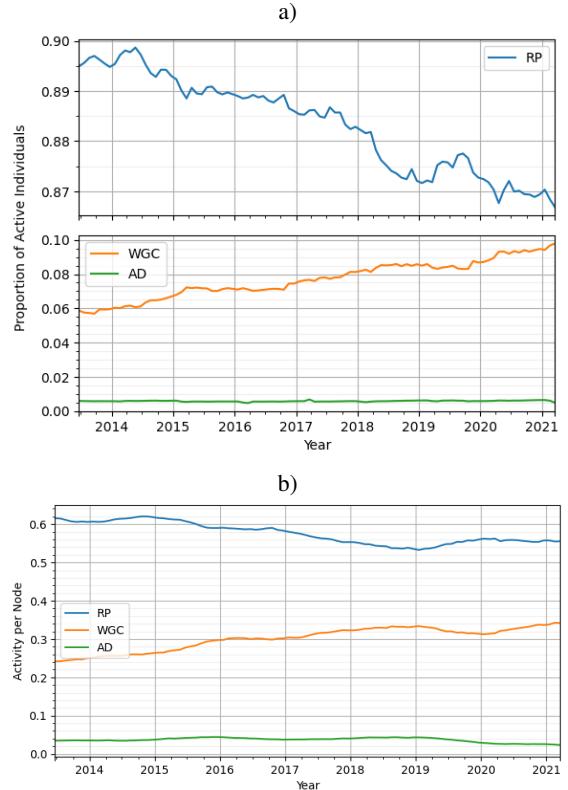


Figure 41: Proportion of regular participants, working group chairs, and area directors by (a) number of active individuals and (b) mailing list activity.

| Status | Sent | Received |
|------------|---------------------|---------------------|
| Before WGC | 67.26 ± 99.80 | 70.11 ± 103.37 |
| After WGC | 77.20 ± 102.85 | 85.90 ± 107.38 |
| Before AD | 218.84 ± 226.92 | 230.50 ± 214.75 |
| After AD | 209.21 ± 186.46 | 219.16 ± 189.85 |

Table 6: Mean number of emails received for one year before/after becoming a WGC (AD) for the first time. The error is one standard deviation.

throughout the period. There are about 15 area directors split amongst 15 roles in 7-8 areas with typically 2 area directors per Area. The set of IETF Areas changes slowly, while working groups are created and closed relatively frequently.

Activity prior to appointment: Table 6 shows the amount of email sent and received for working group chairs and area directors on mailing lists in the year before they first take that role, and the year after. Chairs and area directors are both active in using the mailing lists, area directors much more so. However, they don't become more active to any major degree after taking the roles. This may suggest that the individuals who take on higher hierarchy level roles are already contributing at a higher level.

Summary The answer to the question of whether the higher levels of the hierarchy growing or becoming more centralised, and whether they're associated with an increased domination of the

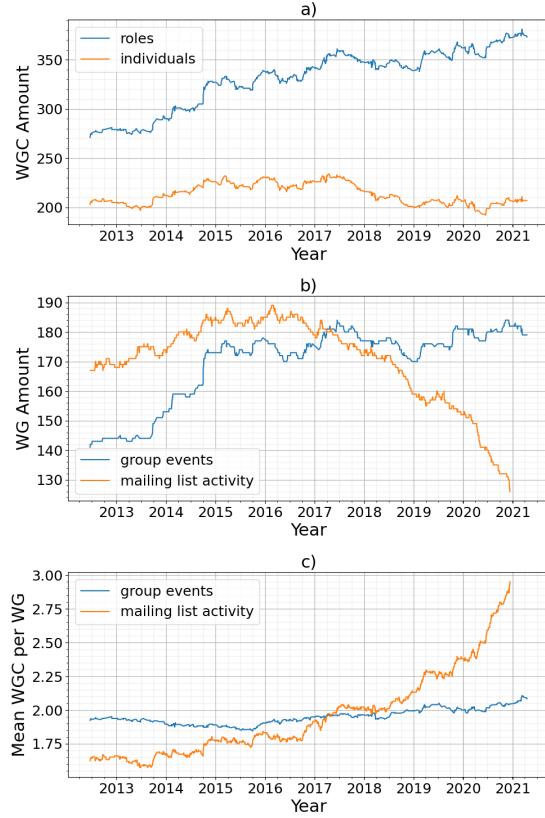


Figure 42: Changes in the number of working groups and working group chairs over time: (a) shows the number of chairs and the number of individuals who are chairs; (b) shows two different estimates of the number of working groups; and c) shows the mean number of chairs per working group. Notice the non-zero-based y-axis exaggerates the variation in these figures.

conversation, is mixed. The number of individuals with working group chair roles has remained broadly constant but the number of such roles has increased, indicating that individuals willing to become chairs have taken on more such roles. The number of the higher level area director roles has, by design, remained broadly constant over the time studied. The proportion of active individuals with working group chairs roles has grown from 6% to 10% and the proportion of the conversation taken by those roles has grown from 25% to 35% over the period studied. Therefore the hierarchy has become more diffused, which may mean regular participants experience more difficulty in expressing their voice.

5.3.2 Organisational Hierarchy and Communication Patterns

In this section, we explore communication patterns within the IETF, to understand how participants in different roles (regular participants, working group chairs, area directors) communicate and whether those patterns of communication are indicative of healthy cooperation in the running of the organisation [60, 30].

For instance, if those in leadership positions send lots of email while receiving little, this suggests a lack of receptiveness of the

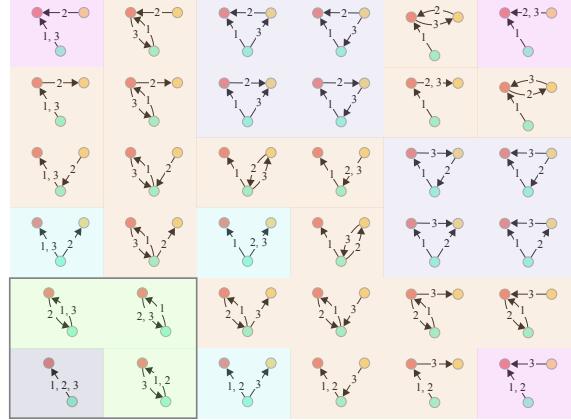


Figure 43: Classification of three edge motif types with the numbers representing the order of communication. Light blue squares are “Outward Star” motifs; purple are “Inward Star” motifs; orange are “Mixed Star” motifs and grey are “Triangle” motifs. The box in the bottom left hand corner represents motifs with only two nodes.

leadership to the suggestions from regular participants, and low confidence of those regular participants to voice their opinions. In contrast, if the IETF operates as outlined in its mission statement, in a more cooperative manner, then communication between those in different roles will be more symmetric.

Temporal Motifs: We compare the ways that people in different roles communicate, looking at the tendency for communications to be purely inbound, purely outbound, or to follow more complex patterns. We describe these in terms of *temporal motifs*: patterns in the sequence in which edges are added to the communication graph over time. Specifically, we analyse every possible combination of three temporally ordered links [71] between at most three nodes. The time order in which the nodes occur is important. The direction of edge 1 is arbitrary, but every edge temporally after it has their direction oriented with respect to it.

Figure 43 shows all possible types of interaction and classifies them into five types: motifs with two nodes; “Outward Star”, announcements or dissemination of information; “Inward Star”, questions or condensing of information; “Mixed Star”, one-on-one discussion; and “Triangle” is group discussion with no individual as the focus. Motifs with only two nodes are rare in this network and we do not present the results.

Raphtory [88] is used to count the prevalence of these motifs. We start with the list of 36 different combinations of three directed edges shown in figure 43. The graph is then split into time windows and then the motifs are counted combinatorially for each window; each of a node’s edges are followed from their origin to three edge steps away. Finally, we combine the motif counts using the categories shown in figure 43 by the colour of the squares.

Impact of Role on Communication Patterns: To quantify the effect being in a leadership position has on network wide communication, we calculate the proportion of three node temporal motifs, lasting at most a month, in which each node participates over a year period. These temporal motifs are counted, categorised into hierarchy levels (regular participant, working group chair, area

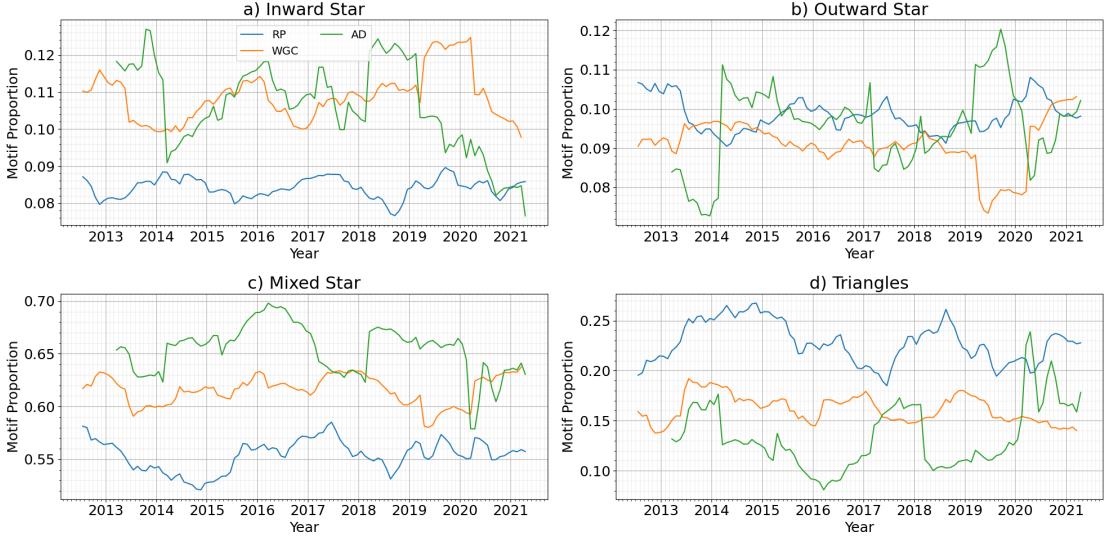


Figure 44: Communication patterns analysed by three edge temporal motifs for RPs,

director) and then proportions are taken of each temporal motif.

Figure 44 shows the proportions of each level for three node “Inward Star”, “Outward Star”, “Mixed Star”, and “Triangles”. The Inward and Mixed Star motifs show about a 4% and 10% increase in proportion for higher levels in the hierarchy versus regular participants. However, about a 10% increase for regular participants is seen for Triangles versus working group chairs and area directors, and Outward Star proportions remain similar for all levels of the hierarchy.

The larger proportion of those in leadership roles in Inward and Mixed Star temporal motifs, whilst Outward Star motifs remain similar for all levels, suggests that those in leadership receive more direct communication. The higher proportion of Triangle motifs for regular participants suggests their discussion is more of a group activity, whereas working group chairs and area directors have a larger proportion of one-on-one conversations. We therefore interpret the general communication patterns within the IETF as a discussion amongst participants interspersed with questions sent to, and announcements from, the working group chairs and area directors. This indicates collaboration between layers and confidence of regular participants’ to voice their opinions upwards.

Figure 45 seems to corroborate this. We calculate the proportion of mailing list threads which nodes in each hierarchy level originate, using a year window, pushed forward by a month each calculation. We see that working group chairs and area directors send a disproportionate number of originating emails to the mailing lists (working group chairs send 30–50% whereas they are 6–10% of active individuals, and area directors 5–10% versus 0.5%). Therefore, the larger proportion of inward motifs for those in leadership roles is in part due to their disproportionately originating email threads. Other IETF participants will then reply to the thread, boosting the proportion of inward motifs working group chairs and area directors appear in. This may suggest the higher levels are good “condensers of communication” as they receive more than they send.

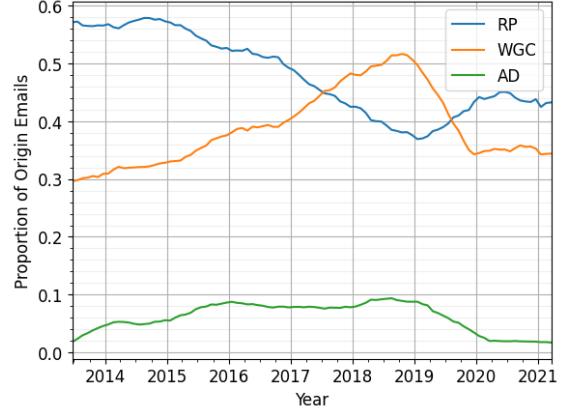


Figure 45: Proportion of “origin” (non reply) emails sent by regular participants, working group chairs, and area directors, this should be interpreted in conjunction with Figure 41 showing the proportion of active individuals in each class.

Summary: Combining the interpretations of Figures 44 and 45, the discussion seems to happen in the following way: area directors or working group chairs are commonly the originators of discussion threads. Where conversation includes a working group chair or area director the conversation tends to be directed around them and usually inward towards the working group chair or area director. The regular participants are more likely to engage in discussions amongst themselves (triangular communication patterns) when higher level individuals are not present.

5.3.3 Communication Up and Down the Hierarchy

Our third hypothesis is that the cooperative nature of the IETF will be evident in the effect that power distance [60] has on communication. The proportion of the communication between

hierarchy levels helps us to understand the confidence of lower levels to exercise their voice: that is, the willingness of regular participants to challenge the working group chairs and of the working group chairs to challenge area directors. This proportion, in a cooperative and voluntary organisation, should be higher for those in lower positions in the hierarchy, showing a confidence in their voice.

We determine the direction communication tends to flow through the hierarchy by looking at inbound and outbound communication flows. The network's edges are categorised based on their source and target node's hierarchy level. For each of the inter-level communication combinations the proportion of communication in the last year going in the “upward” hierarchy direction is calculated. For instance, if there are six emails from regular participants to working group chairs, and twelve from working group chairs to regular participants, then the proportion upward is $\frac{6}{6+12} = \frac{1}{3}$.

In Figure 46, these proportions are plotted, and the year long window is again pushed forward by one month to increase the resolution of changes. A proportion of higher than 0.5 shows more communication flowing up the hierarchy than down. All three levels are mostly upward in direction, with the lowest two levels (regular participant → working group chair) having the largest skew in the proportion. The other two lines (regular participant → area director; working group chair → area director) show similar periodicity, with the working group chairs →AD line dropping as far below 0.5 as above showing a large shift in inter-level communication patterns higher up the hierarchy.

The direct measurement of message flow between the hierarchy levels seems to bolster our interpretation that the cooperative design of the IETF is working. This analysis shows that regular participants are communicating much more preferentially upwards in the hierarchy, despite the decrease in overall share of their activity. The increase within the working group chair level does not have any noticeable effect on their level of communication, if anything there is a negative relationship. This again suggests both working group chairs and area directors perform a kind of “condenser of communication” or “facilitator” role. In other words, this analysis suggests the IETF has a “bottom-up” communication style, where the higher levels encourage the lower levels’ voice. This is the desired pattern of communication for a volunteer organisation that is looking to grow and encourage participants to engage and gradually take on leadership roles.

Summary: The pattern observed was communication flowing up the hierarchy, broadly matching that observed in §5.3.2. Working group chairs and area directors were more likely to receive communication with individuals than send communication to individuals whereas both were more likely than regular participants to send out messages to the list in general.

5.3.4 Mobility Within the Hierarchy

An important question for any organisation is whether higher level participants encourage those at lower levels in the hierarchy to properly engage [60, 39]. The IETF is cooperative and voluntary by design, therefore we hypothesise a high level of encouragement from higher levels to lower. We measure whether activity in one



Figure 46: Proportion of communications that are “up” the hierarchy for different groups. A proportion above 0.5 indicates that the majority of communications are “upward”.

time period is associated with activity in a subsequent time period both for individuals and for neighbourhoods. If those who communicate with higher levels receive a boost in their communication as a result, this is an indication of such encouragement.

Mobility Taxonomy: We aim to measure the association of the activity of nodes, and their neighbours, with their activity in a subsequent time period. To determine this association we use the temporal graph analysis technique called the Mobility Taxonomy [7], which involves correlating the degree and average neighbourhood degree (ND) between two adjoining time windows, both half of the chosen time window length for other analysis. The temporal graph analysis tool Raphtry [88] is used to calculate the raw degree and ND numbers before performing these correlations.

There are six combinations to correlate node degree and ND in time window one and two, the combinations are outlined in detail in [7]. The four measures used in this paper are *Mobility*, *Neighbour Mobility*, *Philanthropy* and *Community*. Each represents the correlation of the degree of a node (or set of nodes) at an earlier or later period of time. *Mobility* can be thought of as the tendency for a node that is active in the first time window to be equally as active in the second. *Neighbour Mobility* is similar but for a node’s neighbourhood. *Philanthropy* can be thought of as the tendency of a node’s neighbourhood to be active in window two if that node is active in window one. Finally, *Community* is the reverse of *Philanthropy* as the activity of a node in window two is compared to that node’s neighbours activity in window one.

Only the nodes which are active in the first window are considered for the second window. Nodes inactive in the second time window are said to have a degree of 0. The ND in the second time window is taken over the same set of neighbours that existed in the first; that is we look at a node’s neighbours in the first window and measure their degree in the second for consistency of comparison.

Participant Mobility: To determine how individuals and their neighbours affect each other’s activity levels in subsequent time periods, we take correlations of degree between different time windows using this Mobility Taxonomy. Figure 47, the time window chosen is one year which is split into two snapshots graphs of six

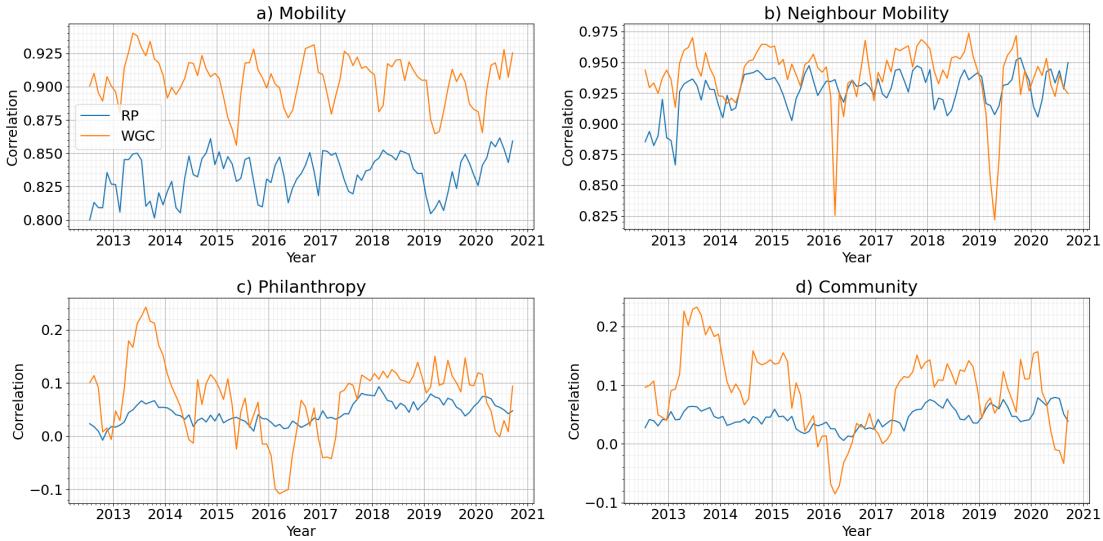


Figure 47: Connection between roles and activity for regular participants and working group chairs viewed in terms of a) Mobility b) Neighbour Mobility c) Philanthropy and d) Community over time.

months, the time window then moves forward one month and the process is repeated. The mid point of the year time window is plotted. Working group chairs and regular participants are plotted but the small number of area directors mean that the correlations are too noisy to show meaningful effects.

The *Mobility* of regular participants and working group chairs is plotted in Figure 47(a). The high amount of correlation between the degree of an individual in the first half of the time window with the second half shows that all users have a tendency to become more active in subsequent time periods if they are highly active in the previous period, but this is more pronounced in working group chairs. This is found to follow a period of about one year for both regular participants and working group chairs, suggesting that a person is likely to be as active as there were a year ago. The *Neighbour Mobility* plot, Figure 47(b), shows a similar tendency for neighbourhoods, but with less of a split.

The *Philanthropy* plot, Figure 47(c), shows a higher correlation for working group chairs than regular participants between an individual's degree in the first snapshot with and their ND in the second. This means that individuals who directly interact with working group chairs become more active on mailing lists subsequently. The *Community* plot, Figure 47(d), shows the correlation between the ND in the first snapshot with the degree in the second. The higher correlation for working group chairs suggests they benefit from interactions with highly active neighbours by themselves becoming more active in subsequent time periods. These trends for *Philanthropy* and *Community* show a reversal in early 2016 that coincides with a period where the mean number of chairs per working group, and the number of working group chairs, rapidly rising, see Figure 42. A merging of the Applications area and Real-time Applications & Infrastructure area in May 2015 may be a cause of this reversal of correlation, as the assignment of working group chairs to areas was changed which may cause more non-reciprocal communication than normal.

Summary: Increased *Mobility* for working group chairs suggests that they gain advantage from their role with their activity leading to similar communication activity in the future. Moreover, they experience an increased effect of *Philanthropy* and *Community*. The former suggests that if a working group chair is engaged in a higher amount of discussion, then in the future the individuals who they were communicating with will also be highly active. The latter suggests that when working group chairs are surrounded by people who are active in discussion, they are likely to be more active themselves in the future too. All of this has less of an effect for regular participants who show a lower correlation in terms of all three aspects. Our interpretation of this is that active working group chairs are “facilitating” discussion in the WG mailing lists; both having their activity boosted by their neighbours and, in turn, boosting the activities of their neighbours. This aligns well with the hypothesis that the cooperative and voluntary design of the IETF lends itself well to higher levels encouraging those lower to engage in discussion.

Working group chairs have a positive effect on their immediate contacts. Those individuals that directly discuss with working group chairs are more likely to engage more in discussion in subsequent time periods. Conversely though the working group chairs who discuss topics with individuals who are active in discussion are, themselves, more likely to engage more fully in discussion. This points to something like a virtuous cycle of chairs encouraging and being encouraged by their direct contacts.

5.3.5 Summary

These findings presented in §5.3.2 – §5.3.4 suggest that the negative association of power distance is not as pronounced for the IETF as the literature finds for traditional commercial organisations. In fact, the facilitation by working group chairs of regular participants suggests the high levels actively encourage lower level participation

in discussion. Moreover, combining this with the finding that the organisational hierarchy has a diffused structure (§5.3.1), with the potential problems this may cause for lower levels, suggests that collaboration between hierarchy levels is high within the IETF.

As the IETF is voluntary and collaborative in its mission, this suggests that their process for tackling the problems caused by a diffused hierarchy and power distance have worked well. One suggested improvement for the IETF is to encourage working group chairs to get involved in more group discussion in the mailing lists. This may help to boost their level of *Philanthropy*, *Community* and *Triangle* motifs, which are indicators of healthy discussion. However, we do not advocate for our analysis techniques to be the only metrics maximised for. Moreover, these conclusions come from analysis of only communication structure not content. A future look into content may elucidate differing relationships between hierarchy levels.

6 Language and Influence

Moving on from discussing the people involved in the IETF (§4) and their interactions (§5), we now consider communications within the IETF from the perspective of language to explore whether there are differences in linguistic patterns and language use in more influential participants, and to consider how the use of language reflects the consensus-driven nature of the standards development process.

Linguistic Inquiry and Word Count (LIWC) [74] is a well-recognised psycholinguistic lexicon; it provides word counts for 85 different linguistic, psychological, personal concern, and informal language marker categories. We generate a *LIWC representation* of the emails sent by a participant in the analysis that follows. This is calculated by aggregating the word counts within each linguistic category for each participant using the LIWC 2015 dictionary (academic license). We filter out 104 ambiguous words that are present in LIWC but have technology, security, and network context meaning in IETF, using manually curated lists, for e.g., attack, argument, secure etc.,

We also normalise by the total number of emails sent by that participant. Such a normalisation is more appropriate here than normalising by total number of words written, as many IETF emails include long technical sections. This generates a representation of a participant as their mean usage of each LIWC category; while this is a relatively reduced, low-dimensional representation of a person's language, it has the advantage of being interpretable and psychologically well-motivated.

A limitation is that we used the standard LIWC-based analysis approach, which is purely lexical and does not take into account the context in which a word appears. Consequently, many words that have very specific senses in the context of the IETF get miscounted as occurrences of LIWC categories. This could be addressed in future by a more advanced method of mapping to LIWC categories that would account for context.

Language and Influence: How do linguistic traits differ between more and less influential participants? That is, are there differences in the way influential participants, based on centrality in the email

social graph, make use of language compared to those who are less influential?

We study mailing list data from the period 2015-2019, comprising 300,806 emails from 5,363 unique participants. We used a 5-year subset of the data due to the computation cost, while still giving a reasonable period to observe the participation consistency in the IETF community, and chose the last period prior to the distortion of the pandemic. For each participant, we calculate the centrality-based influence score and LIWC representation. We fit a linear regression model using LIWC representations to predict influence percentile and observe the magnitude and directions of significant coefficients.

The following LIWC categories are highly correlated ($p < 0.05$) with centrality-based influence: WE, INFORMAL, RISK, ADJECTIVE, ANGER, THEY, and BIO. Categories such as NETSPEAK, SEXUAL, HEALTH, DEATH, BODY are correlated with lower influence. This suggests that influential people tend to indicate a collaborative and community-oriented approach with first-person plural (WE) and third-person plural category (THEY) usage. This is consistent with [48] and [38], who show that influential people use more first-person plural. They also use more organisational language, which is shown by the negative correlation of informal slang language categories (NETSPEAK, SEXUAL, BODY). We see some unexpected hidden trends due to word ambiguity (e.g., words like 'trust' and 'live'), which are discussed below.

Language and Organisational Role: How do linguistic traits vary for participants at different levels of the organisation hierarchy?

Centrality-based influence considers participants based on their email network, but role-based influence, based on the position of a person in the organisational hierarchy (regular participant, working group chair, area director) is equally crucial as they are involved in organisational decision making.¹¹ For the same time period, 2025-2019, we split the data into two categories: (a) emails sent by working group chairs, and (b) emails sent by regular participants who have never been a working group chair (the number of area directors is small, so we exclude them from this analysis). We calculate the LIWC representations for each person, train a logistic regression model to predict category, and observe the LIWC category coefficients.

From 7, we see that working group chairs are more social and collaborative, as is shown by WE and SOCIAL categories. This is in line with our findings for influence-based centrality above and results on leadership engagement from the literature [89, 61, 48, 38]. Chairs also use more tentative statements (TENTAT) in discussions, primarily focused on technical feedback and revisions, or suggesting alternatives. For example: "With the risk of disturbing with statements, but avoiding too many questions: This seems against the goal of reducing headers." and "Question is do we need to carry around an outer IP-in-IP header for that or not?".

Changes in Language with Changes in Role: Finally, we look at how linguistic behaviour of participants changes as they gain influence in the IETF. We consider at participants who went from low to high influence over time: individuals who had a *centrality-based* influence below the 50th percentile when they joined the

¹¹In the top 10% mail-based influential participants, less than 30% are WG chairs with significant role-based influence.

| | | |
|------------------------|-----------------------------------|--|
| Language and influence | High influence | BIO, WE, INFORMAL, THEY, NEGEMO, ANGER, RISK, ADJECTIVE |
| | Low influence | SEXUAL, DEATH, INGEST, NETSPEAK, HEALTH, FEMALE, BODY, AFFILIATION, CONJ |
| Language and role | WG Chair influence | TENTAT, IPRON, SOCIAL, SEE, FEEL, WE |
| | non-WG Chair | COGPROC, RELATIV, AFFILIATION, I, REWARD |
| Changes in Language | Top 10 percentile | ADVERB, PREP, ANGER, AUXVERB, MALE, COGPROC, ACHIEV, RISK, FOCUSPRESENT |
| | Below 50 th percentile | FUNCTION, PPRON, SHEHE, IPRON, NUMBER, CERTAIN, SEXUAL, INFORMAL |

Table 7: LIWC categories where $p < 0.05$.

IETF, and reached the top 10th percentile at some point.

For each participant, we generate two different representations based on two periods—the year of joining, and year of reaching the top 10th percentile for the first time—and assign these to two different classes. We then train a logistic regression model to predict these classes, and examine the coefficients of the LIWC categories

From Table 7, we observe that when participants become influential, according to the centrality-based metric, they are likely to be more descriptive and engaged in immediate state of issues and situations as seen from the correlation of auxiliary verbs (AUXVERB), adverb, risk, and present focus (FOCUSPRESENT). They are also more involved in cognitive processes (COGPROC) as compared to their previous self when they were new to IETF and had little influence.

What Kinds of Language are Used? To better understand these LIWC categories and what kind of words play a role in the behaviour of individual categories, we calculate the frequency of words in each LIWC category as they appear in the emails. Next, we consider the top 30 most frequent words in each LIWC category and perform regression analysis on centrality-based influence for participants, but using only these 30 words as features to generate the participant representation. We conducted this experiment separately for each LIWC category that was significant in the first experiment.

From the word based analysis we make multiple observations. Firstly, words like ‘we’ imply a collective approach and is strongly correlated with the higher influence. Similarly, the use of word ‘well’ is standard, such as politely resuming the conversation (e.g., ‘well, I agree’) or providing an approval over something (e.g., ‘this works as well’). These words are well associated with the influential participants. Otherwise, influential participants are generally not observed to be informal and other frequent words (other than ‘well’) within INFORMAL category do not demonstrate a strong correlation with the growing influence. Also, ‘well’ is the most frequent word in the INFORMAL category.

More influential people, according to both centrality- and role-based metrics, are also observed to engage more with the community. Conversations can often reflect situations where, as a part of review and feedback process, more influential people highlight limitations in protocol standards, stress on specifics, and compare with existing protocols or previous versions. Several words across different LIWC categories (RISK, NEGEMO, and ADJ) highlight such behaviour, e.g., ‘problems’, ‘before’, ‘particular’, ‘specific’, ‘different’, ‘most’, and ‘than’.

However, there are many words with dual sense, like ‘trust’ which has a very technology specific usage related to network

security instead of conversations involving trust issues between individuals or trust in any given situation. Similarly, the word ‘live’ is related with an application or network being live, instead of its conventional meaning. We also observed that some of the LIWC categories, such as BIO, did not have specific terms that could clearly establish its significance in favour of influential participants (e.g., word ‘problems’ and ‘trust’ reflecting the significance for the category RISK), instead such categories had several words with quite weak correlation with influential participants. Such words collectively drifted the weight of the category towards influential participants.

Use of Sensitive Language: We briefly consider the use of sensitive language, relating to politics, race, and religion, in IETF email communications.

Using email data for 2019, the last year prior to the pandemic, we categorise the emails according to the role of their sender, based on the relative social power of the sender within a working group: *working group chair* for emails sent by the current chair of the working group to which the email was sent; *other chair* for emails sent by the chair of a different working group; and *regular participant* for emails sent by people who were not, and had not been at the time of sampling, a workgroup chair. Due to their small numbers, we exclude area directors.

We use three LIWC categories to infer the presence of potentially sensitive language in the emails: POLITICS—words commonly used in political discussions; ETHNICITY—words that identify national, regional, linguistic, ethnic, or racial identities; and RELIGION—use of religious words. Each email is coded according to the working group and role of the sender. The emails are pre-processed to remove subject lines, headers, and any embedded quotes from previous emails, yielding 101,857 texts by 2,300 individuals across 176 group mailing lists. Each text is assigned a score on each of the three language categories using the LIWC software and the resulting scores are averaged for each individual sender for statistical analysis.

Results are shown in Figure 48. The consistent pattern is that people with organizational responsibilities (working group chair, other chair) are less likely to use words potentially connected with sensitive topics than regular participants without organizational responsibilities, even though they are all part of the same discussions on the same email lists. Averaging across the two working group chair categories, the means indicate that people in positions of social power are 3.2 times less likely to use sensitive language overall than those who are not (by category: POLITICS 2.0 times; RELIGION 2.3 times; ETHNICITY 5.5 times). It is worth noting, however, that in absolute terms potentially sensitive language use is

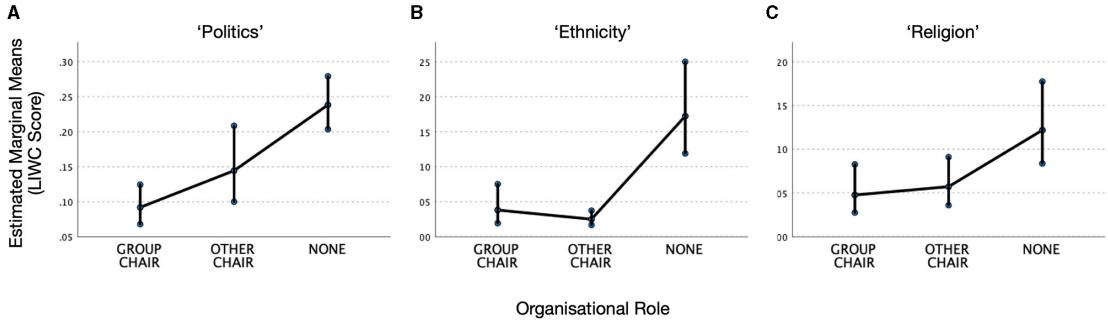


Figure 48: Patterns of language use by organizational role. Estimated marginal means for the LIWC categories: (A) Politics, (B) Ethnicity, and (C) Religion. Error bars show 95

rare in this dataset, as might be expected for primarily work-related communication. The modal LIWC score in both raw data and the averages scores is zero

Summary: Using two aspects of influence, based on centrality in the email social graph and based on organisation role, we were able to unfold several traits that differentiate influential participants from others. Many of our findings seem corroborated by studies in organisational theory. We observed that influential people exhibit more collaborative and community-oriented traits, show stronger signs of engagement in discussions, and are systematically less likely to use potentially sensitive language. We also observed that as people go on to become influential participants, they evolve in their communication and are seen to be more engaging and descriptive in their linguistic style.

7 Success Factors

Finally, we consider factors that lead to success, exploring what makes a successful participant (§7.1) and what makes a document likely to be successfully published as an RFC (§7.2).

Results show that influential authors (i.e., those with high centrality in the email social graph) and authors affiliated with organisations that participate strongly in IETF are more likely to have their documents adopted and published as RFCs. However, documents that build on existing work, have well-defined requirements, limited scope, and meet a need have improved chances of success with authorships playing less of a role. Good, well-connected, people are important, but so is focussed technical work.

7.1 What Makes a Successful Author?

We have identified that influential authors write more drafts (§5.2.3), but are these drafts more likely to see success? That is, does the presence of a well-connected, influential, author make it more likely that an Internet-draft is adopted by a working group for review and development¹² and eventually published as a standards track RFC. Using information from the IETF Datatracker we compile a

¹²There are occasional document, known as Area Director sponsored drafts, that are published as standards track RFCs without going through the working group process, but these too rare (~3% of RFCs) to significantly affect our analysis.

dataset for the *adoption* of a draft by a working group, comprising 11,632 drafts, and employ logistic regression to analyse factors that affect working group adoption of draft.

Methodology: We compute a set of features that may impact adoption or publication for each draft, most of them motivated by the characteristics of influential participants, including the following feature groups: (1) *text*, we derive a term frequency-inverse document frequency (TF-IDF) [84] weighted bag of words vector representation of all email text of messages that either explicitly mention a draft or are part of a thread where the subject line explicitly mentions a draft; (2) *communication patterns*, i.e., volumes of incoming/outgoing communication between draft authors and IETF participants of varying experience levels; (3) Betweenness centrality for authors; (4) *number of emails* sent by authors; (5) *number of drafts* submitted by authors; (6) *length of participation* duration of authors; (7) *number of authors*; (8) *number of IETF areas* in which the author is active; (9) *proportion of top authors* in the top 10 and top 20 percent of influential participants; (10) *topic entropy* (η) scores for authors; (11) *author affiliations* for each of the 20 most influential organizations (20 boolean features); (12) *number of mailing lists* that each author participates in, (13) *number of years active*.

For numerical features that relate to each author individually instead of the draft as a whole (e.g., *number of emails sent*, *centrality*, or *number of years active*, the feature defined for each author separately) in three variants: (i) for the least influential author, (ii) for the most influential author, and (iii) as the average value across all authors of the draft. We compute individual author features using information from the 5-year period prior to draft submission.

Regarding the text based features, we found that the models tend to assign high weights to words such as surnames of active contributors, names of prolific working groups, and common technical terms. While this does make the models perform slightly better, such terms should not be relevant for solving the task as they do not model the relevant part of the conversation. To explore how models would behave in a scenario wherein such words were not available, we construct a domain-specific stop word list, consisting of: (i) one all last names from the Datatracker, (ii) all working group names from the Datatracker, and (iii) technological jargon

| | ¬TXT | | | | TXT | | | |
|--------------------------|------|------|------|------|------|------|------|------|
| | AUC | F1 | P | R | AUC | F1 | P | R |
| Baseline | .500 | .445 | .401 | .500 | .500 | .445 | .401 | .500 |
| Comm. patterns | .644 | .583 | .566 | .601 | .723 | .626 | .609 | .644 |
| Centrality | .655 | .577 | .573 | .581 | .717 | .639 | .625 | .653 |
| Email count | .632 | .566 | .555 | .577 | .718 | .650 | .639 | .662 |
| N. years active | .608 | .569 | .555 | .584 | .702 | .610 | .594 | .626 |
| Number of authors | .578 | .536 | .529 | .544 | .697 | .631 | .621 | .641 |
| Proportion in top | .663 | .594 | .580 | .609 | .728 | .634 | .614 | .655 |
| Draft count | .500 | .445 | .401 | .500 | .696 | .630 | .624 | .636 |
| N. Areas | .634 | .572 | .557 | .588 | .721 | .639 | .622 | .657 |
| N. Mailing lists | .634 | .570 | .557 | .584 | .718 | .641 | .630 | .653 |
| Affiliations | .605 | .584 | .576 | .592 | .708 | .626 | .609 | .643 |
| Topic entropy (η) | .622 | .563 | .550 | .577 | .715 | .642 | .628 | .657 |
| Text | - | - | - | - | .681 | .624 | .622 | .627 |
| Text (S) | - | - | - | - | .667 | .598 | .593 | .614 |
| All feats | .692 | .624 | .600 | .651 | .744 | .645 | .625 | .666 |
| All feats (S) | .692 | .624 | .600 | .651 | .726 | .636 | .627 | .664 |

Table 8: Results for predicting adoption. Each row presents the scores of a model using the corresponding feature group either alone ($\neg TXT$) or combined with the Text features (TXT). *All feats* denotes all features except the text. Rows labeled with (S) use the *Text (S)* variant of the text features.

terms obtained from the web¹³. We remove from this list any terms appearing in top 5,000 English terms (e.g., sometimes people’s names can have identical surface forms as some of the terms relevant to organisational activities). We will refer to this variant as *text (S)*.

To gain empirical insights into which features provide better prediction results, models are run using a feature group alone variant and a variant combining each feature group with the *text* features. This approach was motivated by the finding that *text*, even though conceptually quite different from the rest of the feature set and not our main focus, does provide surprisingly strong results on its own. Thus, we wanted to empirically investigate in more detail how well it complements the rest of the graph-based features. Moreover, we compare all models to a baseline which simply assigns all examples to the majority class.

We split the data into *training* (70%), *development* (10%), and *test* (20%) subsets. As a scoring function, we use the F1 score macro averaged across the positive and negative classes (but we also report area under the curve (AUC), precision, and recall). We train the models on the training set, optimise hyper-parameters on the development set, and report final scores on the test set. The final scores are those obtained by the model variant that fared best on the development set. We use logistic regression implemented in *scikit-learn* [72], given that it is widely used and well interpretable. While we do not perform explicit feature selection, we do have implicit feature selection through L1 regularisation. We consider the regularisation strength a hyper-parameter and consider values from $[2^{-7}, 2^{-6}, \dots, 2^6]$ on the development set. As our data set is quite imbalanced (in a roughly 4:1 ratio, 17% is

adopted), we use different class weights during optimization to counteract this. We employ a non-parametric random shuffling test [100] to check statistical significance of score differences against the baseline. Table 8 summarises prediction results, which will be discussed in the next section.

Next, we wanted to confirm these results at the level of individual features. To this end, we inspect the statistically significant coefficients learned by the model corresponding to each individual feature. For this experiment, we use a slightly different setup. We do not consider the *text* features¹⁴. We begin by first applying the Variance Inflation Factor (VIF) to exclude all features with $VIF > 5$. This, to an extent, mitigates the collinearity that we know exists as some features are by construction highly correlated. We then standardise each remaining feature and then fit a Logistic Regression model from the *statsmodels* package [86] on the entire data set. Table 9 presents the results of this experiment.

Predicting Adoption: To confirm empirical implications from Section 5.2.3, and provide a more in-depth look at the features, we perform statistical analysis summarised in Table 9. The most interesting insight is that the coefficients corresponding to the *Centrality of the most influential author* and *Proportion of authors in the top 10 percentile in influence*, are, indeed, positive and among the largest in absolute value. This highlights that among the people who are influential in the email networks, there is a substantial number of individuals who are proficient in contributing to successful drafts.

Another interesting observation is that author affiliations (which are also indirectly connected to influential participants, as demon-

¹³<https://www.computerhope.com/jargon.htm>, we used a union of the *Internet terms*, *network terms*, and *security terms* categories.

¹⁴There are thousands of text features (terms), making it prohibitively complex to include them. Moreover, they are not central to our research questions about the IETF community.

| Feature | Weight | P-value |
|---------------------------------|---------|---------|
| Centrality of most influential | 0.1691 | 0.000 |
| Centrality of least influential | -0.0698 | 0.048 |
| Proportion of authors in top 10 | 0.1520 | 0.000 |
| Has a Cisco author | 0.1301 | 0.000 |
| Has an AT&T author | 0.1003 | 0.000 |
| Has a Juniper author | 0.0738 | 0.000 |
| Has a China Mobile author | -0.0697 | 0.000 |
| Has an Alcatel-Lucent author | 0.0474 | 0.022 |
| Topic entropy max | 0.1728 | 0.000 |
| Topic entropy min | -0.1135 | 0.002 |

Table 9: Weights of the model that are statistically significant at $p \leq 0.05$ grouped by similarity of features – author influence measures vs. affiliation features vs. topic entropy (η).

strated in §5.2.3, Figure 39) are an important feature group, particularly affiliations such as *Cisco* and *AT&T*. Moreover, topic entropy is also an important feature group, which was also shown to be related to influential participants in our analysis from §5.2.3, Figure 35.

Summary: The statistical analysis shows that influence of draft authors in the email networks does impact the possibility of a draft getting adopted (Tables 8 & 9). This might hint at the ability of participants who hold a domain expertise to be able to engage better with the community. Several WG chairs are already in the top percentile influential category in both the email and co-authorship networks before taking up these leadership roles, which further elevates after taking up such leadership roles (Figure 38). Our analysis also shows that being affiliated with a prominent organisation positively impacts the chances of a draft in getting adopted by a working group, thereby directly driving the innovation process (Table 8 & 9).

7.2 What Makes a Successful Document?

The trends that we observed in Section 3 highlight that, over time, the complexity of the standardisation process is increasing: RFCs take longer to produce, and relate to a larger number of other documents; authors come from a changing and diversifying set of affiliations and countries; and the volume of email interaction for each RFC is increasing, with differences in how senior and junior authors and contributors interact. We next seek to understand how these factors might impact the successful deployment of an RFC in-the-wild, and to test if such success can be predicted.

We build on the work of Nikkhah et al. [68]. However, while the focus of that paper was on a set of features that were derived from the standards documents themselves, we augment this with a set of features derived from the author and email interaction datasets described in the previous section.

To explore the factors that influence protocol deployment, we follow a multi-step process:

1. We reproduce the logistic regression model in [68], to obtain similar results (as shown in Table 10). That paper published

| Model | F_1 | AUC | $F_{1\text{macro}}$ |
|-------------------------------|-------|------|---------------------|
| Most frequent class | .757 | .500 | .379 |
| Baseline | .758 | .616 | .597 |
| Baseline + FS | .762 | .650 | .610 |
| Most frequent class | .724 | .500 | .379 |
| Baseline | .670 | .559 | .547 |
| Baseline + FS | .690 | .620 | .563 |
| Logistic regression all feats | .728 | .724 | .666 |
| Logistic reg. all feats + FS | .820 | .822 | .789 |
| Decision tree all feats + FS | .822 | .838 | .788 |

Table 10: Classifier scores on the entire dataset (251 RFCs, above) and those with our features available (155 RFCs), with or without feature selection (FS).

an expert annotated dataset, labelling RFCs as “successfully deployed” or not. This pertains to whether the RFC was implemented in-the-wild and saw widespread (largely commercial) uptake. This dataset covers 251 RFCs published between 1983 and 2011, and includes 20 features. We create a baseline logistic regression model with $F_1 = 0.762$, AUC = 0.650. The results are similar to those reported in [68] (AUC = 0.670).

2. We curate a set of additional *classification features* driven by our findings in Section 3. These are derived from the RFCs, their authors, and the mailing list interactions. Given that these features are only available for RFCs where Datatracker metadata is available, we model these features across a subset of the original labelled RFCs from [68], looking only at the 155 RFCs where all features can be calculated. We will show that modelling these RFCs with the expanded feature set substantially expands the predictive power of the (Step 1) baseline logistic regression model, with $F_1 = 0.820$, AUC = 0.822.
3. We then train classification models to predict if an RFC will be deployed, using our expanded feature set. We test the model with the manually labelled RFCs and show that we can successfully classify RFCs as deployed or not, with a decision tree-based model having $F_1 = 0.822$, AUC = 0.838.

Classification Features: We begin with the set of features that were originally derived in [68]: IETF Area; Scope (Local, End-to-End (E2E), Bounded (BN), or Unbounded (UB)); Type (New (N), New with Incumbent (NI), Backward compatible extension (EB), or Extension (E)); Change to others (CO); Scalability (SCAL); Security (SCRT); Performance (PERF); Adds value (AV); and Network effect (NE). We refer the reader to [68] for a full explanation of each of these feature.

We further define an additional set of document-based features, derived from our characterisations in §3: days from first draft to RFC publication; number of drafts before RFC publication; number of outbound citations to RFCs/Internet-Drafts; page count; number of inbound citations from articles listed in Microsoft Academic, one and two years after publication; number of inbound citations from other RFCs published within one and two years after

publication; RFC updates or obsoletes a previous RFC; keywords per page; and topics (where we use Latent Dirichlet Allocation (LDA) [13] to induce 50 topics on the texts of all existing RFCs, and use the 50-dimensional probability distribution over topics for a given RFC as the feature vector).

Next, we derive a set of author-based features, from the work described in §4: number of authors on the RFC; if at least one author of previously published RFC; if at least one author in North America, Europe, or Asia; if at least one author from Cisco, Huawei, or Ericsson; authors with diverse affiliations; authors in more than one continent; at least one academic author; and at least one consultant author.

Finally, we derive a set of features based on email interactions, from §5.1: number of emails mentioning the Internet-Drafts that precede publication of the RFC; mean number of emails sent to all RFC authors, for each sender contribution-duration category (young, mid-age, senior); mean number of contributors sending emails to any RFC author, for each contribution-duration category; number of emails sent to most junior and senior RFC authors, for each sender contribution-duration category; and number of contributors sending emails to most junior and senior RFC authors, for each contribution-duration category.

Modelling Methodology: With our expanded feature set, we next build a classifier to predict the success (i.e., deployment) of RFCs.

Our expanded feature set, including all variants of each feature, is very large, totalling 177 features. Given that our additional features can only be calculated for a limited number of data points (155 RFCs), it would be infeasible to conduct statistical analysis, and to build a predictive model, as there are too few data points. To address this, we take a number of steps to reduce the feature space, while maintaining interpretability:

1. Since the largest feature groups are the topics (50) and interaction features (54) we reduce both by applying the χ^2 test to leave only the top 5 features in each group.
2. We remove collinearity by using the Variance Inflation Criterion (VIF), removing all features with $VIF > 5$.
3. We apply forward Feature Selection (FS) to identify features of high predictive value, following [68]. Starting from an empty feature set, in each iteration of the forward procedure, we expand the feature set with the feature that provides the largest increase in the *AUC* score. The procedure ends when there are no more unused features that yield score improvements over the current feature set.

The final set of features with and without FS is given in Tables 11 and 12 respectively.

Using the above feature set, we then train two classification models relying on logistic regression and a decision tree. For assessing predictive performance of the models we use leave-one-out cross-validation, while for the final statistical analysis we fit a logistic regression model on the entire dataset and report the statistically significant coefficients (at significance level $p \leq 0.1$).

Evaluative Results: Using the trained models, we next evaluate their accuracy and explore which features are most predictive of RFC deployment.

| Feature Name | Coef. | $P_i - Z$ |
|--|----------------|--------------|
| Change to others (CO) | 0.0001 | 1.000 |
| Adds value (AV) | 0.7828 | 0.009 |
| Security (SCRT) | 0.3830 | 0.253 |
| Scalability (SCAL) | 0.8755 | 0.100 |
| Performance (PERF) | 0.5108 | 0.323 |
| Microsoft Academic citations, 1 year | 0.2380 | 0.234 |
| Updates others (Yes) | 0.2877 | 0.514 |
| Obsoletes others (Yes) | 1.5315 | 0.001 |
| Keywords per page | 0.3409 | 0.083 |
| Inbound RFC citations, 1 year | 0.6112 | 0.011 |
| Junior-author → Senior (messages) | 0.1226 | 0.463 |
| Young → Senior-author (messages) | 0.2168 | 0.244 |
| Senior → Senior-author (people) | -0.2902 | 0.096 |
| -00 draft mentions | -0.2198 | 0.187 |
| Final draft mentions | 0.1469 | 0.370 |
| All draft mentions (normalised) | -0.0504 | 0.755 |
| -00 draft mentions (normalised) | -0.1052 | 0.525 |
| Author count | -0.0941 | 0.561 |
| Days to publication | 0.1560 | 0.340 |
| Draft Count (DC) | 0.1845 | 0.262 |
| Outbound citation count | 0.2256 | 0.173 |
| Page count | 0.3468 | 0.054 |
| Topic 13 (MPLS) | -0.5629 | 0.068 |
| Topic 19 | -1.2596 | 0.111 |
| Topic 31 | -2.0698 | 0.021 |
| Topic 44 | 0.3992 | 0.129 |
| Topic 45 | 0.3289 | 0.097 |
| Area (INT) | -0.1671 | 0.683 |
| Area (OPS) | 0.5108 | 0.323 |
| Area (SEC) | 0.2231 | 0.638 |
| Area (TSV) | 0.5108 | 0.323 |
| Type, Backward Compatible (EB) | 0.3502 | 0.135 |
| No incumbent | 0.6061 | 0.039 |
| Has incumbent | -0.2007 | 0.655 |
| Scope, End-to-end (E2E) | 0.5878 | 0.035 |
| Scope, Local (L) | 1.3863 | 0.215 |
| Scope, Unbounded (UB) | -1.0986 | 0.033 |
| Has author in N. America (Unknown) | -0.6061 | 0.232 |
| Has author in Europe (Unknown) | 0.1671 | 0.683 |
| Has author in Asia (Yes) | -0.8755 | 0.100 |
| Has author from Cisco (Unknown) | -0.4055 | 0.442 |
| Has author from Cisco (Yes) | 0.4463 | 0.163 |
| Has author from Huawei (Yes) | -1.3863 | 0.215 |
| Has author from Ericsson (Unknown) | 0.0001 | 1.000 |
| Has author from Ericsson (Yes) | 0.5596 | 0.372 |
| Has continent diversity (Yes) | -0.1911 | 0.538 |
| Has an academic author (Yes) | -0.0870 | 0.835 |
| Has a consultant author (Yes) | -0.6931 | 0.423 |

Table 11: Logistic regression w/o feature selection. Statistically significant rows ($p \leq 0.1$) are highlighted.

As evaluative metrics, we use the F_1 score and the area under the ROC curve (*AUC*), as in [68]. We find the standard F_1 score

| Feature Name | Coef. | P _i —Z— |
|--|----------------|--------------------|
| Inbound RFC citations, one year | 0.611 | 0.011 |
| Obsoletes others (Yes) | 0.531 | 0.001 |
| -00 draft mentions | -0.219 | 0.187 |
| Scope, Unbounded (UB) | -1.0986 | 0.033 |
| Has author in Asia (Yes) | -0.8755 | 0.100 |
| Has author in N. America (Unknown) | -0.6061 | 0.232 |
| Keywords per page | 0.3409 | 0.083 |
| Topic 45 | 0.3289 | 0.097 |
| Has an academic author (Yes) | -0.087 | 0.835 |
| Type, Backward Compatible (EB) | 0.3502 | 0.135 |
| Topic 13 (MPLS) | -0.5629 | 0.068 |
| Topic 31 | -2.0698 | 0.021 |
| Topic 44 | 0.3992 | 0.129 |
| Has continent diversity (Yes) | -0.1911 | 0.538 |
| Young → Senior-author (messages) | 0.2168 | 0.244 |
| Topic 19 | -1.2596 | 0.111 |
| Draft Count (DC) | 0.1845 | 0.262 |
| Scope, Local (L) | 1.3863 | 0.215 |
| Has author from Ericsson (Yes) | 0.5596 | 0.372 |

Table 12: Statistical analysis using logistic regression w/ feature selection. Statistically significant rows ($p \leq 0.1$) are highlighted.

gives overly optimistic performance estimates due the data being skewed towards the positive class. Consequently, we also report an $F1_{macro}$ score that takes into account both the positive and negative class, reflecting performance more realistically. As a baseline comparison, we present results of our re-implementation of [68] using their feature set.

The prediction results of the models and feature sets are summarized in Table 10. Our baseline model performs at a similar level to that from [68]. We further obtain considerable performance improvements with our additional features. Furthermore, we find that additional performance can be obtained by using feature selection. The best performing model, with an F1 score of 0.822, is the decision tree trained on the entire feature set. Note, we also tested several non-linear models (neural networks, support vector machines with non-linear kernels). These attained similar or worse results as our decision tree model, and therefore we omit them due to space constraints.

We next investigate the most important features in predicting success. We posit that these can offer useful insight for working group chairs and RFC editors. The feature importance results are presented in Tables 11 and 12. Some meaningful patterns emerge here. We observe features such as adding value, being scalable, obsoleting other RFCs, page count, increased keyword usage, no incumbent RFC to compete against and limited scope all being positively correlated with the likelihood of an RFC becoming deployed. In contrast, we find that having a broad (unbounded) scope negatively impacts deployment.

We also observe some curious trends that speak to the limitations of our dataset. The model finds that having an author in Asia is negatively correlated with deployment, though with only borderline statistical significance. Further analysis shows that only 10% of

labelled RFCs have an author in Asia. As shown in §4, we see that demographics of the IETF are changing, with a recent notable increase in author representation from Asia. This might suggest that new authors need time to learn what makes a deployable protocol; equally our deployment data may be biased towards the Internet in North America and Europe. This finding requires much more exploration.

We further see that some of the topics extracted by LDA are also useful. For example, Topic 13 is characterised by a cluster of terms associated with MPLS, a widely deployed routing protocol. We find a negative correlation with this topic, but this is likely because there is a significant number of RFCs that propose modifications or additional features for the protocol, some of which do not see deployment. Overall, most of the results are in line with expectations, but more annotated data is needed for further insights.

Discussion: Features associated with building on existing RFCs positively correlate with deployment. RFCs that obsolete earlier versions of the same protocol are likely to be deployed, indicating that the IETF community is good at identifying and maintaining protocols that are seeing use. This is also highlighted by the significance of the inbound RFC citation and “adds value” features. These show that RFCs that are later cited by other documents, or that add value to other protocols in the stack, are more likely to be deployed. Protocols that are extensible, and can be extended and adapted for new uses, are more likely to see deployment. The discussion in [25] on careful choice of extension points resonates; as do the implications for starting new work.

We find that having end-to-end scope (i.e., where only the endpoints of a connection needs to implement the RFC) is positively correlated with deployment, whereas having an unbounded scope (i.e., where the entire Internet may need to be updated) is negatively correlated. This indicates that well-scoped RFCs, that are cognisant of their deployment challenges, are more likely to see deployment. A recent IAB workshop [5] noted that deployment often occurs in unforeseen ways: limiting the changes needed to deploy a protocol facilities this.

The number of keywords (e.g., SHOULD, MUST) used per page, which shows the number of normative requirements an RFC imposes on implementations, also correlates with deployment. This suggests that well-specified requirements lead to robust, interoperable, implementations that see wide use. Unsurprisingly, features such as citation rates are also predictive of deployment.

Notwithstanding the limitations of the dataset, we find that the majority of author demographic features are not significant. While §4 identified recent shifts in the demographics of authors, these do not appear to have a major influence in the deployment of the RFCs that they write.

Our results broadly support the conclusions of [92], where the IAB noted that, in addition to being technically sound, a successful protocol must meet a need, be incrementally deployable, and be open to extension and maintenance.

Summary: Our results show that RFCs that build on existing work, have well-defined requirements, limited scope, and meet a need have improved chances of being widely deployed. These document-based features have implications for how the IETF should design and specify protocols.

We found that the majority of our author-based features, including geographic and affiliation diversity, did not significantly impact the deployment prospects of RFCs. Further work is needed to expand the dataset to ensure that this result holds across a larger set of RFCs.

Finally, the focus of our modelling in this section has been on the impact of document, author, and interaction-based features on protocol deployment. However, this is only one part of the lifecycle of an RFC. It remains to consider the impact of these, and other, features on the key stages of an Internet-Draft’s development towards becoming an RFC, such as working group adoption.

8 Related Work

Documents, Demographics, Communication Patterns: To the best of our knowledge, were the first to perform a comprehensive statistical analysis of IETF activities considering trends in document publication (§3), author demographics (§4), and communication patterns. §5.1. That said, there have been several prior efforts that have focused on specific aspects of the IETF. Arkko, a former Chair of the IETF, maintains a website (<https://www.arkko.com/tools/stats.html>) that provides various statistics about the IETF, including about its documents, authors and their affiliations. The DataTracker also provides some statistics on documents and authors (<https://datatracker.ietf.org/stats/>). Attendance statistics are reported in the IETF plenaries (available in the proceedings at <https://www.ietf.org/meeting/past/>).

BigBang [9] is a Python toolkit for analysing online collaborative communities through mailing list data. Niedermayer et al. [67] discuss the challenges of working with large mailing list datasets, including conducting entity resolution where contributors have changeable pseudonyms. Our work addresses a number of these challenges by augmenting the mailing list data with metadata from the IETF DataTracker.

Huitema [46] carried out an evaluation of a small subset of RFCs, to understand the various sources of delay in their publication. The author observed that the main source of delay was the working group process, highlighting the need to better understand the dynamics of that process. In addition, the author found that citation counts did not correlate well with the adoption or deployment of specific RFCs.

Influence and Impact: Our work to characterise participant influence and impact (§5.2) builds on a rich literature regarding influence in online social networks and other communities including studies of Twitter [23, 98, 3], Instagram [101] and decentralised social networks [10, 41]. For example, [99] measured the impact of follower influence on message propagation. Others have focused on devising metrics to capture influence: similarly to [56, 98, 87], we use betweenness centrality as a metric of influence.

Online collaborative communities such as Wikimedia and Open Stack have also been studied for understanding communities, collaboration, strategy making, and organisational structures, through mailing lists and platform interactions [29, 15, 102]. Bosou and Carver [16] highlight that contributors with much higher reputation

are successful in seeking reviews from the community in a much shorter time span and are more likely to get their suggestions accepted. This is well aligned with most of our findings related to the impact of influential people on the Internet-Draft adoption process. Zhang et al. observe [102] that in the open source software (OSS) ecosystem, several participating organisations (firms or companies) may engage in intentional or passive collaborations, or they may contribute in an isolated way. They find that an organisation’s influence in the collaboration network is positively correlated with its scale of contribution within the ecosystem. In some ways, this is similar to our observation related to prominent affiliated organisations in the co-authorship network and their ability to collaborate and produce drafts.

Building social interaction networks out of email communication data is a well trodden field [8, 64], and the introduction of time evolution into the analysis [70, 97] is recent but has quickly gained interest. However, the analysis of email based communication is less prevalent with the notable exception of the well-known, and oft analysed, Enron emails [54]. The communication dataset we examine is interesting for several reasons: it is a long duration email dataset, it is annotated with meta data representing participant roles, and it represents discussion in a decentralised consensus driven organisation [81, 1] rather than a business with a traditional management hierarchy. While the Enron data set is often used for classification of emails into categories [62], there are also several studies of organisation interaction patterns on it. Namely, [90, 55] are concerned with developing methodologies for studying how communication evolves over time. In contrast, [28] focus on providing insights more than on developing a methodology, revealing that organisational structure is reflected in email communication patterns. Other studies focusing on mailing list data include OSS mailing lists [11, 82].

Temporal Patterns in Communication: Analysing the temporal patterns in communication, as we do for the IETF organisational hierarchy in §5.3, is an emerging field of research [63]. Prior work has analysed the time evolution of important nodes in a network [34], whether their neighbours are similarly important [73] and some others the correlations of the trajectory of snapshots of networks over time [58]. Compared to this, we brought new analysis techniques, that expand on the idea of tracking the evolution of node importance, to an email list corpus.

A form of temporal network analysis that is key to understanding communication in social networks is computing temporal motifs [71, 57] between nodes. This technique involves look at the sequence in which edges attach to nodes over time which lets us see how conversation happens with fine resolution. Applications include understanding the diffusion of information through a social network using undirected motifs [85], and analysing the collaboration and scientific mobility of co-authorship networks [14]. To the best of our knowledge, temporal motifs have not yet been used for analysis of communication in social networks which means knowledge of high time-resolution direct communication between individuals is ripe for the picking.

The effects of hierarchy on the “voice” of individuals within an organisation is a key aspect [75] of Organisational Behaviour (OB) research. A steep [47] hierarchy can have a negative effect

on voice by reducing the variety of perspectives [2] in higher levels, or positive effect where decisions can be made quickly. However, a diffused hierarchy may also have negative effects [47] on voice by contributing to lower level “bystanders” who let the higher levels do the talking, or a positive effect with a variety of approaches to management. Similarly, a large “power distance” can cause problems to lower level voice if the higher levels are “authoritarian” [30], however if they are “benevolent” [60, 39] then power distance is less of a problem. We tested these ideas in regard to the voluntary and collaborative nature of the IETF organisation, whereas the literature mostly focuses on traditional organisations.

Use of Language: Considering use of language (§6), most studies of influence either focus on community structure rather than language, or use language indirectly. Urena et al. [93] give a survey of the former approach. In an example of the latter, Prabhakaran [79] compares users with different influence in terms of their linguistic similarity or co-adaptation, the increasing similarity of interlocutors to each other in how they use language (see also [27, 96, 69, 51, 20]). Some studies focus on modelling influence from text of Enron emails by identifying keywords/phrases that indicate influence [19, 37]. Rosenthal [83] and Tchokni et al. [91] extend this approach to other domains, including Twitter, Wikipedia talk pages, and debates, and include a wider range of linguistic markers.

Success Factors: The work in §7 builds, in part, on Nikkhah et al. [68] by statistically exploring RFC adoption. We extended this work with improved modelling, and by incorporating additional features that characterise not only the RFCs themselves but also the standardisation process that produced them.

9 Conclusions

We explored the key standardisation activities within the IETF, one of the primary technical standards developing organisations for the Internet, considering trends in document production, complexity, and correctness (§3); the people involved in Internet standards development and how the demographics of the community have shifted over time (§4); interaction between participants, who has influence and impact in the IETF and how that influence is evident in their activities, and the dynamics of communications within the IETF (§5); use of language, to explore whether there are differences in linguistic patterns and language use in more influential participants, and to consider how the use of language reflects the consensus-driven nature of the standards development process (§6); and factors that lead to success, exploring what makes a document likely to be published as an RFC and what makes a successful participant (§7).

In the process, we have revealed a vibrant ecosystem, albeit one with a number of key challenges:

1. The rate of standards development has slowed over time and the IETF is showing signs of struggling with managing the complexity of standards development.
2. While authorship is becoming more diverse as the IETF grows, there are still significant issues with representation

from certain regions, including Africa and South America. If the IETF is to become more representative, then further effort is clearly required to build more diverse engagement, both in terms of geography, organisational affiliation, and gender diversity [49].

3. Despite significant growth over time, a relatively small subset of well-connected and highly influential participants remain an important driving force in the IETF. The community is growing better connected, and has become less fragmented, over time, and is growing in resilience to departure of influential participants, but there is more to be done.

On the positive side, we observed that influential people in IETF exhibit collaborative and community-oriented traits, and also stronger signs of engagement in discussions. We also observed that as people go on to become influential participants, they evolve in their communication and are seen to be more engaging and descriptive in their linguistic style and patterns of communication. The IETF has a reputation for being a difficult community to engage with (“loud men, talking loudly”) that is not altogether unwarranted [22, 21], but there is also evidence that those in positions of influence, and leadership roles, are broadly cooperative and supportive of newcomers and those with less experience, within their own cultural framework.

Further research is needed to fully understand the IETF and its role in the Internet governance ecosystem, but our analysis to date shows a complex and evolving organisation—with some expected growing pains—that is meeting the ongoing challenges of developing technical standards for the modern, global, Internet.

Acknowledgements

This work has been supported in part by the UK Engineering and Physical Sciences Research Council under grants EP/S033564/1 and EP/S036075/1 (“Streamlining Social Decision Making for Improved Internet Standards”).

References

- [1] H. T. Alvestrand. A Mission Statement for the IETF. RFC 3935, Oct. 2004.
- [2] C. Anderson and C. E. Brown. The functions and dysfunctions of hierarchy. *Research in organizational behavior*, 30:55–89, 2010.
- [3] I. Anger and C. Kittl. Measuring influence on twitter. In *Proceedings of the 11th international conference on knowledge management and knowledge technologies*, pages 1–4, Graz, Austria, 2011. ACM.
- [4] J. Arkko. IETF blog: Diversity. <https://www.ietf.org/blog/diversity>, Apr. 2013. Accessed: 2021/09/27.
- [5] J. Arkko and T. Hardie. Report from the IAB workshop on design expectations vs. deployment reality in protocol development. RFC Editor, Feb. 2021. RFC 8980.

- [6] M. Barnes, M. Karan, S. McQuistin, C. S. Perkins, G. Tyson, I. Castro, R. G. Clegg, and M. Purver. Temporal network analysis of email communication patterns in a long standing hierarchy. In *Proceedings of the International Conference on Web and Social Media (ICWSM)*, Buffalo, NY, USA, June 2024. AAAI.
- [7] M. R. Barnes, V. Nicosia, and R. G. Clegg. Measuring equality and hierarchical mobility on abstract complex networks. In *Conference on Complex Networks (Complenet)*, pages 15–28. Springer, 2023.
- [8] F. Belanger. Communication patterns in distributed work groups: A network analysis. *IEEE Transactions on Professional Communication*, 42(4):261–275, 1999.
- [9] S. Benthall. Testing generative models of online collaboration with bigbang (pp. 182–189). In *Proceedings of the 14th Python in Science Conference*, 2015.
- [10] H. Bin Zia, A. Raman, I. Castro, I. Anaobi, E. D. Cristofaro, N. Sastry, and G. Tyson. Toxicity in the decentralized web and the potential for model sharing. In *Proceedings of ACM SIGMETRICS*, 2022.
- [11] C. Bird, A. Gourley, P. Devanbu, M. Gertz, and A. Swaminathan. Mining email social networks. In *Proceedings of the 2006 international workshop on Mining software repositories*, pages 137–143, 2006.
- [12] H. Birkholz, C. Vigano, and C. Bormann. Concise Data Definition Language (CDDL): A Notational Convention to Express Concise Binary Object Representation (CBOR) and JSON Data Structures. RFC 8610, June 2019.
- [13] D. M. Blei, A. Y. Ng, and M. I. Jordan. Latent dirichlet allocation. *the Journal of machine Learning research*, 3:993–1022, 2003.
- [14] H. D. Boekhout, V. A. Traag, and F. W. Takes. Investigating scientific mobility in co-authorship networks using multilayer temporal motifs. *Network Science*, 9(3):354–386, 2021.
- [15] A. Bosu and J. C. Carver. How do social interaction networks influence peer impressions formation? a case study. In *IFIP International Conference on Open Source Systems*, pages 31–40. Springer, 2014.
- [16] A. Bosu and J. C. Carver. Impact of developer reputation on code review outcomes in oss projects: An empirical investigation. In *Proceedings of the 8th ACM/IEEE international symposium on empirical software engineering and measurement*, pages 1–10, 2014.
- [17] S. Bradner. The Internet standards process – revision 3. RFC Editor, Oct. 1996. RFC 2026.
- [18] S. Bradner. Key words for use in RFCs to indicate requirement levels. RFC Editor, Mar. 1997. RFC 2119.
- [19] P. Bramsen, M. Escobar-Molano, A. Patel, and R. Alonso. Extracting social power relationships from natural language. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 773–782, Portland, Oregon, USA, June 2011. Association for Computational Linguistics.
- [20] J. A. Buske. Linguistic accommodation between leaders and followers. B.S. thesis, University of Twente, 2019.
- [21] C. Cath. The technology we choose to create: Human rights advocacy in the Internet Engineering Task Force. *Telecommunications Policy*, 45(6), July 2021.
- [22] C. Cath. Loud men talking loudly: Exclusionary cultures of Internet governance. Technical report, Critical Infrastructure Lab, University of Amsterdam, Apr. 2023.
- [23] M. Cha, H. Haddadi, F. Benevenuto, and K. Gummadi. Measuring user influence in twitter: The million follower fallacy. In *Proceedings of the international conference on web and social media*, page 8, Washington, DC, USA, 2010. Association for the Advancement of Artificial Intelligence.
- [24] D. Chen, L. Lü, M.-S. Shang, Y.-C. Zhang, and T. Zhou. Identifying influential nodes in complex networks. *Physica a: Statistical mechanics and its applications*, 391(4):1777–1787, 2012.
- [25] D. D. Clark, J. Wroclawski, K. R. Sollins, and R. Braden. Tussle in cyberspace: Defining tomorrow’s Internet. In *Proceedings of the SIGCOMM Conference*, Pittsburgh, PA, USA, Aug. 2002. ACM.
- [26] D. Crocker and N. Clark. An IETF with Much Diversity and Professional Conduct. RFC 7704, Nov. 2015.
- [27] C. Danescu-Niculescu-Mizil, L. Lee, B. Pang, and J. Kleinberg. Echoes of power: Language effects and power differences in social interaction. In *Proceedings of the 21st international conference on World Wide Web*, pages 699–708, 2012.
- [28] J. Diesner, T. L. Frantz, and K. M. Carley. Communication networks from the enron email corpus “it’s always about the people. enron is no different”. *Computational & Mathematical Organization Theory*, 11(3):201–228, 2005.
- [29] L. Dobusch and J. Kapeller. Open strategy-making with crowds and communities: Comparing wikimedia and creative commons. *Long Range Planning*, 51(4):561–579, 2018.
- [30] J. Duan, C. Bao, C. Huang, and C. T. Brinsfield. Authoritarian leadership and employee silence in china. *Journal of Management & Organization*, 24(1):62–80, 2018.
- [31] S. Farrell, F. Badiei, B. Schneier, and S. M. Bellovin. Reflections on Ten Years Past the Snowden Revelations. RFC 9446, July 2023.

- [32] S. Farrell and H. Tschofenig. Pervasive Monitoring Is an Attack. RFC 7258, May 2014.
- [33] B. Fenner, M. J. Handley, I. Kouvelas, and H. Holbrook. Protocol Independent Multicast - Sparse Mode (PIM-SM): Protocol Specification (Revised). RFC 4601, Aug. 2006.
- [34] M. Fire and C. Guestrin. The rise and fall of network stars: Analyzing 2.5 million graphs to reveal how high-degree vertices emerge over time. *Information Processing & Management*, 57(2):102041, 2020.
- [35] H. Flanagan. Fifty years of RFCs. RFC Editor, Dec. 2019. RFC 8700.
- [36] Z. Ghalmame, M. El Hassouni, C. Cherifi, and H. Cherifi. Centrality in modular networks. *EPJ Data Science*, 8(1):15, 2019.
- [37] E. Gilbert. Phrases that signal workplace hierarchy. In *Proceedings of the ACM 2012 conference on Computer Supported Cooperative Work*, pages 1037–1046, 2012.
- [38] A. Guinote. How power affects people: Activating, wanting and goal seeking. *Annual review of psychology*, 68:353–381, 2017.
- [39] Y. Guo, Y. Zhu, and L. Zhang. Inclusive leadership, leader identification and employee voice behavior: The moderating role of power distance. *Current Psychology*, pages 1–10, 2020.
- [40] L. Hagen, T. Keller, S. Neely, N. DePaula, and C. Robert-Cooperman. Crisis communications in the age of social media: A network analysis of zika-related tweets. *Social Science Computer Review*, 36(5):523–541, 2018.
- [41] A. I. Hassan, A. Raman, I. Castro, H. B. Zia, E. De Cristofaro, N. Sastry, and G. Tyson. Exploring content moderation in the decentralised web: The pleroma case. In *Proceedings of ACM CoNEXT*, 2021.
- [42] X. He, N. Meghanathan, et al. Correlation of eigenvector centrality to other centrality measures: random, small-world and real-world networks. *NeCoM, CSITEC*, pages 09–18, 2016.
- [43] P. G. T. Healey, P. Khare, I. Castro, G. Tyson, M. Karan, R. Shekhar, S. McQuistin, C. S. Perkins, and M. Purver. Power and vulnerability: Managing sensitive language in organisational communication. In *Proceedings of the Society for Text and Discourse*, Oslo, Norway, June 2023.
- [44] P. G. T. Healey, P. Khare, G. Tyson, M. Karan, I. Castro, R. Shekhar, S. McQuistin, C. S. Perkins, and M. Purver. Power and vulnerability: Managing sensitive language in organisational communication. *Frontiers in Psychology*, 14, Feb. 2024.
- [45] M. Hoffman, F. Bach, and D. Blei. Online learning for latent dirichlet allocation. *advances in neural information processing systems*, 23:856–864, 2010.
- [46] C. Huitema. Evaluation of a sample of RFCs produced in 2018. RFC Editor, Jan. 2021. RFC 8963.
- [47] I. Hussain, R. Shu, S. Tangirala, and S. Ekkirala. The voice bystander effect: How information redundancy inhibits employee voice. *Academy of Management Journal*, 62(3):828–849, 2019.
- [48] E. Kacewicz, J. W. Pennebaker, M. Davis, M. Jeon, and A. C. Graesser. Pronoun use reflects standings in social hierarchies. *Journal of Language and Social Psychology*, 33(2):125–143, 2014.
- [49] M. Kaeo. Experience of women participating in the IETF. Report to IETF LLC, Oct. 2023.
- [50] M. Karan, P. Khare, R. Shekhar, S. McQuistin, C. S. Perkins, I. Castro, G. Tyson, P. G. T. Healey, and M. Purver. LEDA: a large-organization email-based decision-dialogue-act analysis dataset. In *Findings of the Association for Computational Linguistics*, pages 6080–6089, Toronto, ON, Canada, July 2023.
- [51] K. Kawabata, V. Berisha, A. Scaglione, and A. LaCross. A convex model for linguistic influence in group conversations. In *INTERSPEECH*, pages 1442–1446, 2016.
- [52] P. Khare, M. Karan, S. McQuistin, C. S. Perkins, G. Tyson, M. Purver, P. Healey, and I. Castro. The web we weave: Untangling the social graph of the IETF. In *Proceedings of the International Conference on Web and Social Media (ICWSM)*, pages 500–511, Atlanta, GA, USA, June 2022. AAAI.
- [53] P. Khare, R. Shekhar, M. Karan, S. McQuistin, C. S. Perkins, I. Castro, G. Tyson, P. G. T. Healey, and M. Purver. Tracing linguistic markers of influence in a large online organisation. In *Proceedings of the Association for Computational Linguistics*, Toronto, ON, Canada, July 2023.
- [54] B. Klimt and Y. Yang. The enron corpus: A new dataset for email classification research. In *European Conference on Machine Learning*, pages 217–226. Springer, 2004.
- [55] G. Kossinets, J. Kleinberg, and D. Watts. The structure of information pathways in a social communication network. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 435–443, 2008.
- [56] N. Kourtellis, T. Alahakoon, R. Simha, A. Iamnitchi, and R. Tripathi. Identifying high betweenness centrality nodes in large social networks. *Social Network Analysis and Mining*, 3(4):899–914, 2013.
- [57] L. Kovanen, M. Karsai, K. Kaski, J. Kertész, and J. Saramäki. Temporal motifs in time-dependent networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2011(11):P11005, 2011.

- [58] L. Lacasa, J. P. Rodriguez, and V. M. Eguiluz. Correlations of network trajectories. *Physical Review Research*, 4(4):L042008, 2022.
- [59] L. Leslie Daigle and Internet Architecture Board. The rfc series and rfc editor. RFC Editor, July 2007. RFC 4844.
- [60] X. Li and L. Xing. When does benevolent leadership inhibit silence? the joint moderating roles of perceived employee agreement and cultural value orientations. *Journal of Managerial Psychology*, 36(7):562–575, 2021.
- [61] A. H. Liu. Pronoun usage as a measure of power personalization: A general theory with evidence from the chinese-speaking world. *British Journal of Political Science*, 52(3):1258–1275, 2022.
- [62] G. Madjarov, D. Kocev, D. Gjorgjevikj, and S. Džeroski. An extensive experimental comparison of methods for multi-label learning. *Pattern recognition*, 45(9):3084–3104, 2012.
- [63] N. Masuda and R. Lambiotte. *A guide to temporal networks*. World Scientific, 2016.
- [64] M. McPherson, L. Smith-Lovin, and J. M. Cook. Birds of a feather: Homophily in social networks. *Annual review of sociology*, 27(1):415–444, 2001.
- [65] S. McQuistin, M. Karan, P. Khare, C. S. Perkins, M. Purver, P. G. T. Healey, I. Castro, and G. Tyson. Errare humanum est: What do RFC errata say about Internet standards? In *Proceedings of the Network Traffic Measurement and Analysis Conference*, Naples, Italy, June 2023. IFIP.
- [66] S. McQuistin, M. Karan, P. Khare, C. S. Perkins, G. Tyson, M. Purver, P. Healey, W. Iqbal, J. Qadir, and I. Castro. Characterising the IETF through the lens of RFC deployment. In *Proceedings of the Internet Measurement Conference*, Online, Nov. 2021. ACM.
- [67] H. Niedermayer, N. Schwellnus, D. Raumer, E. Cordeiro, and G. Carle. Information mining from public mailing lists: A case study on ietf mailing lists. In *International Conference on Internet Science*, pages 301–309. Springer, 2017.
- [68] M. Nikkhah, A. Mangal, C. Dovrolis, and R. Guérin. A statistical exploration of protocol adoption. *IEEE/ACM Transactions on Networking*, 25(5):2858–2871, 2017.
- [69] B. Noble and R. Fernández. Centre stage: How social network position shapes linguistic coordination. In *Proceedings of the 6th workshop on cognitive modeling and computational linguistics*, pages 29–38, 2015.
- [70] P. Panzarasa, T. Opsahl, and K. M. Carley. Patterns and dynamics of users’ behavior and interaction: Network analysis of an online community. *Journal of the American Society for Information Science and Technology*, 60(5):911–932, 2009.
- [71] A. Paranjape, A. R. Benson, and J. Leskovec. Motifs in temporal networks. In *ACM International Conference on Web Search and Data Mining*, pages 601–610, 2017.
- [72] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [73] N. Pedreschi, D. Battaglia, and A. Barrat. The temporal rich club phenomenon. *Nature Physics*, 18(8):931–938, 2022.
- [74] J. W. Pennebaker, R. L. Boyd, K. Jordan, and K. Blackburn. The development and psychometric properties of LIWC2015. Technical report, 2015.
- [75] J. Pfrombeck, C. Levin, D. D. Rucker, and A. D. Galinsky. The hierarchy of voice framework: the dynamic relationship between employee voice and social hierarchy. *Research in Organizational Behavior*, page 100179, 2023.
- [76] J. Postel. Internet Control Message Protocol. RFC 792, Sept. 1981.
- [77] J. Postel. Internet Protocol. RFC 791, Sept. 1981.
- [78] J. Postel. Transmission Control Protocol. RFC 793, Sept. 1981.
- [79] V. Prabhakaran, A. Arora, and O. Rambow. Staying on topic: An indicator of power in political debates. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1481–1486, 2014.
- [80] R. Řehůřek and P. Sojka. Software Framework for Topic Modelling with Large Corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50, Valletta, Malta, May 2010. ELRA.
- [81] P. Resnick. On Consensus and Humming in the IETF. RFC 7282, June 2014.
- [82] P. C. Rigby and A. E. Hassan. What can oss mailing lists tell us? a preliminary psychometric text analysis of the apache developer mailing list. In *Fourth International Workshop on Mining Software Repositories (MSR’07: ICSE Workshops 2007)*, pages 23–23. IEEE, 2007.
- [83] S. Rosenthal. Detecting influencers in social media discussions. *XRDS: Crossroads, The ACM Magazine for Students*, 21(1):40–45, 2014.
- [84] G. Salton and M. J. McGill. *Introduction to modern information retrieval*. McGraw-Hill, New York, 1983.
- [85] S. Sarkar, R. Guo, and P. Shakarian. Using network motifs to characterize temporal network evolution leading to diffusion inhibition. *Social Network Analysis and Mining*, 9:1–24, 2019.

- [86] S. Seabold and J. Perktold. statsmodels: Econometric and statistical modeling with python. In *9th Python in Science Conference*, 2010.
- [87] A. Solé-Ribalta, M. De Domenico, S. Gómez, and A. Arenas. Centrality rankings in multiplex networks. In *Proceedings of the conference on Web science*, pages 149–155, Bloomington, IA, USA, 2014. ACM.
- [88] B. Steer, F. Cuadrado, and R. Clegg. Raphtory: Streaming analysis of distributed temporal graphs. *Future Generation Computer Systems*, 102:453–464, 2020.
- [89] T. Strzalkowski, S. Shaikh, T. Liu, G. A. Broadwell, J. Stromer-Galley, S. Taylor, U. Boz, V. Ravishankar, and X. Ren. Modeling leadership and influence in multi-party online discourse. In *Proceedings of COLING 2012*, pages 2535–2552, 2012.
- [90] L. Tang, H. Liu, J. Zhang, and Z. Nazeri. Community evolution in dynamic multi-mode networks. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 677–685, 2008.
- [91] S. E. Tchokni, D. O. Séaghdha, and D. Quercia. Emoticons and phrases: Status symbols in social media. In *Eighth International AAAI Conference on Weblogs and Social Media*, 2014.
- [92] D. Thaler and B. Aboba. What makes for a successful protocol? RFC Editor, July 2008. RFC 5218.
- [93] R. Urena, G. Kou, Y. Dong, F. Chiclana, and E. Herrera-Viedma. A review on trust propagation and opinion dynamics in social networks and group decision making frameworks. *Information Sciences*, 478:461–475, 2019.
- [94] US Bureau of Industry and Security, Commerce. Addition of entities to the entity list. Rule 84 FR 22961, May 2019.
- [95] T. W. Valente, K. Coronges, C. Lakon, and E. Costenbader. How correlated are network centrality measures? *Connections (Toronto, Ont.)*, 28(1):16, 2008.
- [96] G. Ver Steeg and A. Galstyan. Information-theoretic measures of influence based on content dynamics. In *Proceedings of the sixth ACM international conference on Web search and data mining*, pages 3–12, 2013.
- [97] B. Viswanath, A. Mislove, M. Cha, and K. P. Gummadi. On the evolution of user interaction in Facebook. In *Workshop on Online Social Networks*, pages 37–42, 2009.
- [98] L. Weitzel, P. Quaresma, and J. P. M. de Oliveira. Measuring node importance on twitter microblogging. In *Proceedings of the 2nd International Conference on Web Intelligence, Mining and Semantics*, pages 1–7, Craiova, Romania, 2012. ACM.
- [99] S. Ye and S. F. Wu. Measuring message propagation and social influence on twitter. com. In *International conference on social informatics*, pages 216–231, Laxenburg, Austria, 2010. Springer.
- [100] A. Yeh. More accurate tests for the statistical significance of result differences. In *Proceedings of the 18th International Conference on Computational Linguistics*, Saarbrücken, Germany, 2000. Association for Computational Linguistics.
- [101] K. Zarei, D. Ibosiola, R. Farahbakhsh, Z. Gilani, K. Garimella, N. Crespi, and G. Tyson. Characterising and detecting sponsored influencer posts on instagram. In *International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*, pages 327–331, Online, 2020. IEEE/ACM.
- [102] Y. Zhang, M. Zhou, K.-J. Stol, J. Wu, and Z. Jin. How do companies collaborate in open source ecosystems? an empirical study of openstack. In *2020 IEEE/ACM 42nd International Conference on Software Engineering (ICSE)*, pages 1196–1208. IEEE, 2020.